

Clustering validation

Is there any real clustering? How good is it?

- Book: Chapter 6.9
- External material: Halkidi et al. (2002): Cluster Validity Methods: Part I. ACM SIGMOD Record 31(2): 40–45.
<https://doi.org/10.1145/565117.565124>

Three similar problems

1. Clustering tendency: is there any clustering in data presented with certain features?
2. Determining number of clusters (or other parameters)
3. Evaluating goodness of clustering
 - compare different methods
 - compare against classification

All three depend on the **clustering objective!**

- assumptions on clusters (e.g., compactness, shape)
- separation between clusters

Evaluating goodness of clustering

1. Internal criteria

- validity indices, similar to objective functions
- do not work, if clustering had a different objective!
- can be used to i) evaluate a single clustering or ii) compare clusterings (as **relative indices**)

2. External criteria

- compare clustering to a predefined classification
- classes may not reflect natural clusters

3. Statistical hypothesis testing

- maybe the most sound approach, but computationally demanding

Internal validity indices

- indices assume some clustering objective → reward methods with the same objective
 - even a good clustering can get a bad score if a different objective!
 - many indices assume/favor spherical or convex clusters
- best for comparing similar algorithms and tuning parameters
- Some popular indices:
 - **Average silhouette**
 - **Calinski-Harabasz index**
 - **Davies-Bouldin index**

Silhouette index

Silhouette of a point \mathbf{x} is

$$S(\mathbf{x}) = \begin{cases} 0 & \text{if } \mathbf{x} \text{ a cluster of its own} \\ \frac{b-a}{\max\{a,b\}} & \text{otherwise} \end{cases}$$

$$a = \text{avg}\{d(\mathbf{x}, \mathbf{y}) \mid \mathbf{x} \in C, \mathbf{y} \in C\}$$

$$b = \min_q \text{avg}\{d(\mathbf{x}, \mathbf{y}) \mid \mathbf{x} \in C, \mathbf{y} \in C_q, C \neq C_q\}$$

\approx how closely \mathbf{x} matches its own cluster and how loosely the neighbouring cluster

- $S(\mathbf{x}) \in [-1, 1]$, **high values good**
- **Average silhouette** describes goodness of entire clustering
- flexible: any distance function d

Example: Silhouette of points


What negative values mean?

$$S(\mathbf{x}) = \begin{cases} 0 & \text{if singleton} \\ \frac{b-a}{\max\{a,b\}} & \text{otherwise} \end{cases}$$

$$a = \text{avg}\{d(\mathbf{x}, \mathbf{y}) \mid \mathbf{x} \in C, \mathbf{y} \in C\}$$

$$b = \min_q \text{avg}\{d(\mathbf{x}, \mathbf{y}) \mid \mathbf{x} \in C, \mathbf{y} \in C_q, C \neq C_q\}$$

image source http://www.sthda.com/english/wiki/wiki.php?id_contents=7952



Calinski-Harabasz index

$$S_{CH} = \frac{(n - K)B}{(K - 1)W}$$

- **between-cluster variance** $B = \sum_{i=1}^K |C_i| L_2^2(\mathbf{c}_i, \mathbf{m})$, where \mathbf{m} is the mean of the whole data
- **within-cluster variance** $W = \sum_{i=1}^K \sum_{\mathbf{x} \in C_i} L_2^2(\mathbf{x}, \mathbf{c}_i)$
- requires $K \geq 2$
- range $[0, \infty[$, **high values good**
- When could you get value 0?

Calinski-Harabasz index (cont'd)

$$S_{CH} = \frac{(n - K)B}{(K - 1)W} = \frac{(n - K) \sum_{i=1}^K |C_i| L_2^2(\mathbf{c}_i, \mathbf{m})}{(K - 1) \sum_{i=1}^K \sum_{\mathbf{x} \in C_i} L_2^2(\mathbf{x}, \mathbf{c}_i)}$$

Note: $W = SSE(\mathbf{C})$. K -means criterion minimizes $W \Rightarrow$ maximizes B , because

$$\sum_{\mathbf{x} \in \mathcal{D}} L_2^2(\mathbf{x}, \mathbf{m}) = \sum_{i=1}^K \sum_{\mathbf{x} \in C_i} L_2^2(\mathbf{x}, \mathbf{c}_i)^2 + \sum_{i=1}^K |C_i| L_2^2(\mathbf{c}_i, \mathbf{m})$$

$\Rightarrow S_{CH}$ favours especially K -means!

Important: need to use L_2 in clustering!

Davies-Bouldin index

$$S_{DB} = \frac{1}{K} \sum_{i=1}^K \max_{j \neq i} \frac{S_i + S_j}{D_{ij}}, \text{ where}$$

- $S_i = \left(\frac{1}{|C_i|} \sum_{\mathbf{x} \in C_i} L_p^q(\mathbf{x}, \mathbf{c}_i) \right)^{\frac{1}{q}}$ measures dispersion of C_i
 - usually $q = 2$ (stdev of distances)
 - if $q = 1$, average distances
- $D_{ij} = L_p(\mathbf{c}_i, \mathbf{c}_j)$ measures separation between C_i and C_j
- max: for each C_i , evaluate relation to most problematic C_j
- possible to take avg instead of max

Important: use the same L_p as the clustering algorithm!

Davies-Bouldin index (cont'd)

$$S_{DB} = \frac{1}{K} \sum_{i=1}^K \max_{j \neq i} \frac{S_i + S_j}{D_{ij}}, \text{ where}$$


$$S_i = \left(\frac{1}{|C_i|} \sum_{\mathbf{x} \in C_i} L_p^q(\mathbf{x}, \mathbf{c}_i) \right)^{\frac{1}{q}} \text{ and } D_{ij} = L_p(\mathbf{c}_i, \mathbf{c}_j)$$

- range $[0, \infty[$, **small values good**
- When could you get value 0?

Possible strategies when S_{DB} used to determine K :

- restrict number of singletons (e.g., 0 or a few)
- define $S_i = a$ for some large a , when $|C_i| = 1$

External validation: Compare clustering against predefined classification



A confusion matrix: clustering vs. classification

	Class 1	Class 2	Class 3	
Cluster 1	n_{11}	n_{12}	n_{13}	m_1
Cluster 2	n_{21}	n_{22}	n_{23}	m_2
Cluster 3	n_{31}	n_{32}	n_{33}	m_3
	c_1	c_2	c_3	n

image source Cunnigham <https://slideplayer.com/slide/14318989/>

External validation

Given clustering C_1, \dots, C_K and classification D_1, \dots, D_q .
Many validation indices! E.g.,

- **purity**

$$Pur(C) = \frac{1}{n} \sum_{i=1}^K \max_j |C_i \cap D_j|$$

- be careful! (increases with K)
- **normalized mutual information NMI** (robust, independent of K)
- **Rand index**

Normalized mutual information

Normalized mutual information by Strehl and Ghosh (2003):

$$NMI = \frac{I(C, D)}{\sqrt{H(C)H(D)}}$$

mutual information $I = \sum_{C_i \in C} \sum_{D_j \in D} P(C_i, D_j) \log \frac{P(C_i, D_j)}{P(C_i)P(D_j)}$


entropy $H(C) = - \sum_{C_i \in C} P(C_i) \log P(C_i)$

- + does not depend on the number of clusters
- many singleton clusters can cause problems

Note: Also other variants of normalized mutual information, give always equation and/or reference what you use!

Statistical hypothesis testing: motivation

SI can be pretty good even for random data!



- each feature generated independently from uniform distribution
- 100 randomizations
- K -means repeated 100 times \rightarrow best result for each K

Experiment by Georgy Ananov for MDM 2023

Statistical hypothesis testing


Procedure:

1. decide a **null hypothesis** H_0 to test
 - describes the state where there isn't any clustering
 - e.g., H_0 : All sets of n locations in certain region are equally likely.
2. decide a **test statistic** T
 - may be a validity index
3. What is the probability to obtain at least as good test statistic values as in data (where $T = t$) if H_0 was true?

Statistical hypothesis testing

Assume that large T value good

Idea: If $P(T \geq t)$ very small
 \Rightarrow unlikely that the observed clustering had occurred by chance



- $P(T \geq t)$ is the **p-value** that can be used as a significance measure

Statistical hypothesis testing

Problem: How to evaluate p -value? (T 's distribution seldom known!)

- often by Monte Carlo experiments (randomization tests):
 - generate random data sets fulfilling H_0 , cluster them and evaluate T
 - p -value \approx proportion of random sets that obtained $T \geq t$ (if large T good)
- computationally demanding (a lot of simulations!)
- many alternatives for H_0 s and T s

Other evaluation: What the clustering reveals?

- Look at cluster sizes (e.g., C_1 : $n - 2$ data points and C_2 : 2 points – likely outliers!)
- How do the clusters differ? (selected and external features)
 - e.g., rats clustered by body measurements (weight, tail and body length, organ weights)
 - 2 clusters: big and small rats
 - vs. 3 clusters: C_1 : young or sick rats, C_2 : pregnant or nursing females, C_3 : other adults
- Are all clusters clear? (e.g., C_1 and C_3 intermingled, C_2 separate)

Summary

- Remember validation, but be cautious!
 - even random data can produce clusterings, but they seldom pass validation
 - problem: indices biased or do not reflect the underlying clustering
 - try always more than one validation technique
- Objective, distance measure, clustering method and validation should match!

Sources and further reading

- Halkidi et al. (2001): On clustering validation techniques, Journal of Intelligent Information Systems 17: 107–145. https://www.researchgate.net/publication/2500099_On_Clustering_Validation_Techniques
- Jain and Dubes (1988): Algorithms for clustering data, Ch 4.
- Gan, Ma, Wu (2007): Data clustering - theory, algorithms, and applications, Ch 17, https://www.researchgate.net/publication/220694937_Data_Clustering_Theory_Algorithms_and_Applications

Sources and further reading

- Vargha, Bergman, Takacs: Performing Cluster Analysis Within a Person-Oriented Context: Some Methods for Evaluating the Quality of Cluster Solutions. *Journal of Person-Oriented Research*, 2: 78-86, 2016.

Spectral clustering

Contents:

- Matrices from the similarity graph
- 1D spectral embedding & clustering
- Unnormalized and normalized spectral clustering
- Important choices

Book: Sections 2.4.4.3, 6.7, 19.3.4


Recommended external material:

von Luxburg (2007): A Tutorial on Spectral Clustering.




Presemo: <https://presemo.aalto.fi/mdm2023>

Recap: How could you cluster these?

Arbitrary shapes?



Only similarity graphs known?



Images: White (2019) <https://www.markhw.com/blog/word-similarity-graphs>, Cooper (2021) <https://spin.atomicobject.com/2021/09/07/spectral-clustering/>, Park & Kim (2020) <https://doi.org/10.1115/DETC2020-22642>, Scikit-learn documentation https://ogrisel.github.io/scikit-learn.org/sklearn-tutorial/auto_examples/cluster/plot_cluster_comparison.html

General idea of graph-based clustering

1. Present data as a similarity (neighbourhood) graph G
 2. Cluster nodes of G with a network clustering or community detection algorithm
- + can detect arbitrary-shaped clusters
 - + even varying cluster densities (given k nearest neighbour similarity graph)
 - + for any data type (if pairwise similarity/distance defined)
 - computationally costly
 - many parameter choices

Spectral clustering: Idea

1. Create similarity graph G

- node v_i for the i th data point ($i = 1, \dots, n$)
- edge weight w_{ij} = similarity between nodes v_i and v_j

2. Present data in (low-dimensional) vector space (i.e., find vectors $\mathbf{y}_1, \dots, \mathbf{y}_n$) such that **local similarity/clustering structure is preserved**

- idea: choose \mathbf{Y} to minimize $cost(\mathbf{G}, \mathbf{Y}) = \sum \sum w_{ij} L_2^2(\mathbf{y}_i, \mathbf{y}_j)$
- intuition: large w_{ij} tends to produce small $d(\mathbf{y}_i, \mathbf{y}_j)$
- \rightarrow easy after reformulation with a Laplacian matrix


3. Cluster \mathbf{y}_i s with K -means (etc.)

What is needed?

From \mathbf{G} derive:

1. weight matrix \mathbf{W}
2. diagonal degree matrix $\mathbf{\Lambda}$
3. Laplacian matrix $\mathbf{L} = \mathbf{\Lambda} - \mathbf{W}$
4. normalized Laplacian matrices \mathbf{L}_{rw} , \mathbf{L}_{sym} if desired)

Similarity graph and weight matrix W



0.00	0.00	1.00	0.00	0.33	0.00
0.00	0.00	0.00	0.33	0.60	0.33
1.00	0.00	0.00	0.00	0.33	0.00
0.00	0.33	0.00	0.00	0.00	0.33
0.33	0.60	0.33	0.00	0.00	0.00
0.00	0.33	0.00	0.33	0.00	0.00

- W adjacency matrix of a weighted graph
- $W_{ij} = w_{ij}$ (similarity between nodes v_i and v_j)
- if unweighted graph, use weights 1 (edge) or 0

Diagonal degree matrix Λ ($\Lambda_{ii} = \sum_{j=1}^n W_{ij}$)

$$\Lambda = \begin{bmatrix} 1.33 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1.26 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1.33 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0.66 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1.26 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0.66 \end{bmatrix}$$

$$W = \begin{bmatrix} 0.00 & 0.00 & 1.00 & 0.00 & 0.33 & 0.00 \\ 0.00 & 0.00 & 0.00 & 0.33 & 0.60 & 0.33 \\ 1.00 & 0.00 & 0.00 & 0.00 & 0.33 & 0.00 \\ 0.00 & 0.33 & 0.00 & 0.00 & 0.00 & 0.33 \\ 0.33 & 0.60 & 0.33 & 0.00 & 0.00 & 0.00 \\ 0.00 & 0.33 & 0.00 & 0.33 & 0.00 & 0.00 \end{bmatrix}$$

(Unnormalized) Laplacian matrix $\mathbf{L} = \mathbf{\Lambda} - \mathbf{W}$

$$\mathbf{L} = \begin{bmatrix} 1.33 & 0.00 & -1.00 & 0.00 & -0.33 & 0.00 \\ 0.00 & 1.26 & 0.00 & -0.33 & -0.60 & -0.33 \\ -1.00 & 0.00 & 1.33 & 0.00 & -0.33 & 0.00 \\ 0.00 & -0.33 & 0.00 & 0.66 & 0.00 & -0.33 \\ -0.33 & -0.60 & -0.33 & 0.00 & 1.26 & 0.00 \\ 0.00 & -0.33 & 0.00 & -0.33 & 0.00 & 0.66 \end{bmatrix}$$

⇒ normalized Laplacian matrices:

- Random-walk Laplacian $\mathbf{L}_{rw} = \mathbf{\Lambda}^{-1}\mathbf{L}$
- Symmetric Laplacian $\mathbf{L}_{sym} = \mathbf{\Lambda}^{-0.5}\mathbf{L}\mathbf{\Lambda}^{-0.5}$

Idea of 1D spectral embedding & clustering

Goal: find embedding $\mathbf{y} = (y_1, \dots, y_n)^T$, where each y_i corresponds v_i and $cost(G, \mathbf{y})$ minimal.

$$cost(G, \mathbf{y}) = \sum \sum w_{ij} (y_i - y_j)^2 = 2\mathbf{y}^T \mathbf{L} \mathbf{y}$$

- we want to avoid trivial solution $\forall i : y_i = 0 \rightarrow$
- scaling constraint (e.g.) $\mathbf{y}^T \mathbf{y} = 1$ (i.e., $\sum_i y_i^2 = 1$)
- \mathbf{L} is positive semidefinite (eigenvalues λ_i real, $\lambda_i \geq 0$)
- solution smallest non-trivial eigenvector of \mathbf{L}

Extra: Why eigenvectors \mathbf{y} of \mathbf{L} would be the solution?

Task: Find \mathbf{y} such that $2\mathbf{y}^T \mathbf{L} \mathbf{y}$ minimal given constraint $\mathbf{y}^T \mathbf{y} = 1$

Method of Lagrange multipliers:

1. Reformulate as a Lagrangian function

$$\mathcal{L}(\mathbf{y}, \lambda) = \mathbf{y}^T \mathbf{L} \mathbf{y} - \lambda(\mathbf{y}^T \mathbf{y} - 1)$$

2. Set the partial derivatives (with respect to \mathbf{y} and λ) as 0
3. Reduces to $\mathbf{L} \mathbf{y} = \lambda \mathbf{y}$ **Eigenvalue & -vector definition!**

Idea of 1D spectral embedding & clustering

- solution smallest non-trivial eigenvector \mathbf{y} of \mathbf{L}
- $cost = 2\mathbf{y}^T \mathbf{L} \mathbf{y} = 2\mathbf{y}^T \lambda \mathbf{y} = 2\lambda(y_1^2 + \dots + y_n^2) = 2\lambda$ (λ eigenvalue)
- $cost$ minimal, when λ minimal (recall $\mathbf{y}^T \mathbf{y} = 1$)
- but **skip trivial solution** $\lambda = 0$ with \mathbf{y} (proportional to)
 $\mathbf{1} = (1, \dots, 1)^T$
 - exists always when \mathbf{G} connected
- \rightarrow optimal solution **eigenvector corresponding to the 2nd smallest λ**
- cluster elements of \mathbf{y} with K -means

Example

Unnormalized Laplacian L

$$\begin{bmatrix} 1.33 & 0.00 & -1.00 & 0.00 & -0.33 & 0.00 \\ 0.00 & 1.26 & 0.00 & -0.33 & -0.60 & -0.33 \\ -1.00 & 0.00 & 1.33 & 0.00 & -0.33 & 0.00 \\ 0.00 & -0.33 & 0.00 & 0.66 & 0.00 & -0.33 \\ -0.33 & -0.60 & -0.33 & 0.00 & 1.26 & 0.00 \\ 0.00 & -0.33 & 0.00 & -0.33 & 0.00 & 0.66 \end{bmatrix}$$

Eigenvalues:

$\approx 0, 0.20, 0.99, 0.99, 1.99, 2.33^*$

Second smallest eigenvector:

$(0.48, -0.19, 0.48, -0.48, 0.19, -0.48)^T$


The new representation can be clustered by K -means:



* 1st eigenvalue $1.9e-16$ due to imprecision (should be 0)


Another example with 1D embedding

Fully connected weighted graph.



Eigenvector:

$(0.32, 0.34, 0.28, 0.34, 0.29, -0.32, -0.27, -0.31, -0.36, -0.31)^T$



$n = 10$

Example by Bruno Ordozgoiti, MDM 2020

Generalization with multidimensional embedding

Unnormalized spectral clustering

Input: Graph G with adjacency matrix W , number of clusters K .

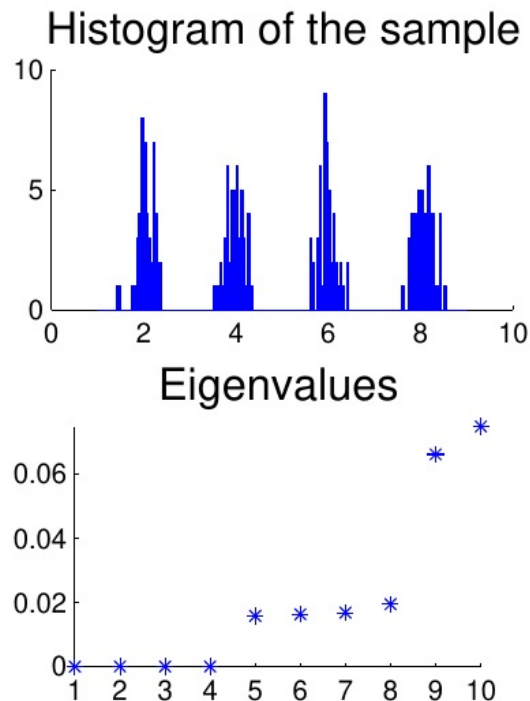
1. Compute the Laplacian $L = \Lambda - W$
2. Compute the **eigenvectors** y_1, \dots, y_k of L corresponding to the k smallest eigenvalues (excluding $\lambda = 0$)
3. Present the data as matrix Y whose columns are y_1, \dots, y_k .
4. Cluster Y with K -means.

Note: Usually $k = K$ or $k < K$. Eigengap $|\lambda_{k+1} - \lambda_k|$ can be used to choose k .



Eigengap heuristic for choosing k

Choose k such that $\lambda_1, \dots, \lambda_k$ small but λ_{k+1} relatively large.

clear gap



no gap



Normalized spectral clustering using random walk

Laplacian \mathbf{L}_{rw}

Input: Graph \mathbf{G} with adjacency matrix \mathbf{W} , number of clusters K .

1. Compute the **random walk** Laplacian $\mathbf{L}_{rw} = \mathbf{\Lambda}^{-1} \mathbf{L}$
2. Compute the right eigenvectors $\mathbf{y}_1, \dots, \mathbf{y}_k$ of \mathbf{L}_{rw} corresponding to the k smallest eigenvalues (excluding $\lambda = 0$)
3. Present the data as matrix \mathbf{Y} whose columns are $\mathbf{y}_1, \dots, \mathbf{y}_k$.
4. Normalize the columns of \mathbf{Y} to unit norm.
5. Cluster \mathbf{Y} with K -means.

Normalized spectral clustering using symmetric normalized Laplacian \mathbf{L}_{sym}

Input: Graph \mathbf{G} with adjacency matrix \mathbf{W} , number of clusters K .

1. Compute the **symmetric normalized Laplacian**
$$\mathbf{L}_{sym} = \mathbf{\Lambda}^{-1/2} \mathbf{L} \mathbf{\Lambda}^{-1/2}$$
2. Compute the eigenvectors $\mathbf{y}_1, \dots, \mathbf{y}_k$ of \mathbf{L}_{sym} corresponding to the k smallest eigenvalues (excluding $\lambda = 0$)
3. Present the data as matrix \mathbf{Y} whose columns are $\mathbf{y}_1, \dots, \mathbf{y}_k$.
4. Normalize the rows of \mathbf{Y} to unit norm.
5. Cluster \mathbf{Y} with K -means.

Important choices


- **Method:** Unnormalized, random walk or symmetric normalized?
 - Usually normalization helps. Suggestion: try random walk first.
- **Similarity measure**
 - should measure local similarity reliably (close neighbours)
 - for numeric data, Gaussian similarity $\exp\left(\frac{-\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$ often used
- **Similarity graph** and its parameters
 - this has a strong effect on results!

Common choices for the similarity graph

General goal: sparse but connected graph (or number of connected components $\ll K$)

1. **ϵ -neighbourhood** graph: keep only $w_{ij} \geq \epsilon$
 - problems if clusters of different densities
2. **k -nearest neighbour** graph: v_i among k nearest neighbours of v_j **or** vice versa
 - often a good first choice
 - can break the graph into disconnected components
3. **mutual k -nearest neighbour** graph: v_i among k nearest neighbours of v_j **and** vice versa

Similarity graph examples (von Luxburg, Fig 3)



Similarity graph (cont)


4. fully connected graph

- often with Gaussian similarity $\kappa(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(\frac{-\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$
(radial basis function, RBF)
- how to choose σ ?
- Note: in scikitlearn parameter $\gamma = \frac{1}{2\sigma^2}$
- graph not sparse \rightarrow heavy computation


Choice of parameters (ϵ, k, σ) affects a lot, too!

Example: neighbourhood with $\kappa(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(\frac{-\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$


Data, query point red




$\sigma = 0.1$




$\sigma = 0.2$



$\sigma = 0.4$



$\sigma = 0.8$



$\sigma = 1.6$





Image source: Bruno Ordozgoiti, MDM 2020 slides

Example: Different clustering results with different similarity graphs



Parameters: $k = 2$, $\epsilon = 0.3$, RBF with $\gamma = 10$ (i.e., $\sigma = \sqrt{5}$)

Experiment by Lai Khoa for MDM 2023

Summary

Idea: similarity graph \rightarrow low-dimensional VS presentation (eigenvectors) \rightarrow clustering (K -means etc.)

- + very powerful (virtually any datatype, arbitrary shapes)
- computationally expensive
 - creating similarity graph $O(n^2)$, spectral decomposition $O(n^3)$
- many important parameter choices

Further reading

von Luxburg: A Tutorial on Spectral Clustering. Statistics and Computing, vol. 17, pp. 395–416, 2007.

Reading guide: Sec 2 overview, 2.2, Sec 3 overview + definitions of Laplacian matrices from 3.1-3.2, Sec 4, Sec 5 overview, (possibly Sec 6 overview), Sec 8. ^a

^asection overview = text before subsections