

Mining Episodes and Episode Rules

Juho Rinta-Paavola

10 October 2023

Who am I?

Juho Rinta-Paavola

- Studied CS at Aalto in 2015–2022
- Did my Master's thesis on episode mining
- Now working as a software engineer for the company I did my thesis for

Overview

- Definitions: what's an episode?
- Overview of approaches to episode mining
- The WINEPI algorithm
- Redundant episode rules
- Practical application of episode mining

Definitions

- **Episode** A is a partially ordered collection of **event types** $L(e) \in \Sigma$ that occur close together in an **event sequence** s
- **Episode rule** $B \rightarrow A$ describes the relationship between an episode A and its **subepisode** $B \preceq A$


Definitions

- **Episode** A is a partially ordered collection of **event types** $L(e) \in \Sigma$ that occur close together in an **event sequence** s
- **Episode rule** $B \rightarrow A$ describes the relationship between an episode A and its **subepisode** $B \preceq A$

Note: much of what you have learned about itemsets and association rules applies with little modification to episodes and episode rules!

Definitions

- Data: exactly one long-ish event sequence



- Event e consists of an event type $L(e)$ and a timestamp $t(e)^*$
- Multiple event sequences can be concatenated, but if you have a lot of short sequences, other kinds of pattern might be more appropriate

* However, the sequence is not required to be a time sequence

Definitions

- What does it mean for events in an episode to occur close together?
- **Window**-constrained: limit maximum duration between first and last event of the occurrence.
- Does the episode occur in the window of size win starting at $t_{\text{win_start}}$, i.e. in the interval between $t_{\text{win_start}}$ and $t_{\text{win_start}} + \text{win}^*$
- Problem: a fixed window size can be simultaneously too long for simple episodes and too short for complex episodes.
- **Gap**-constrained: limit maximum duration between consecutive events in an occurrence.[†]

*Mannila et al., “Discovering Frequent Episodes in Sequences”.

[†]Casas-Garriga, “Discovering Unbounded Episodes in Sequential Data”;
Méger and Rigotti, “Constraint-Based Mining of Episode Rules and Optimal Window Sizes”.

Definitions

- In general, an episode may have multiple copies of an event type, and may require them to be in any partial order
Example: first a occurs, then b and c occur in either order, and finally a occurs again
- If the episode is totally ordered, then it is a **serial** episode
Example: the event types (a, c, b, a) occur in that order
- If the episode is unordered, then it is a **parallel** episode
Example: a occurs twice, b and c occur once each, order doesn't matter
- If the episode has at most one of each event type, then it is an **injective** episode.

Remark: serial and parallel episodes can be thought of as sequences and multisets of event types, respectively

Approaches: frequent episode mining

- Various definitions of episode frequency, e.g.
 - Number of fixed-size windows in s in which episode occurs*
 - Number of non-overlapping **minimal occurrences**[†]
A window of s is a minimal occurrence of A iff A occurs in the window, but not in any smaller window
- Beware! Some definitions of frequency in the literature have later turned out not to be monotone

*Mannila et al., “Discovering Frequent Episodes in Sequences”.

[†]Tatti, “Significance of Episodes Based on Minimal Windows”.

Approaches: statistically significant episodes

- Various null hypotheses are possible
 - If events are independent and identically distributed, then episode A isn't statistically significant ($p > 0.5$) if its non-overlapping occurrence based frequency is $\leq |s|/|A||\Sigma|^*$
 - If events are independent and distributed according to event type frequencies, or according to a Markov model, then distribution of window-based episode frequencies is asymptotically normal[†]
 - If events are independent and distributed according to event type frequencies, then distribution of the sum of lengths of minimal non-overlapping occurrences is asymptotically normal[‡]

*Laxman et al., "Discovering frequent episodes and learning hidden Markov models: a formal connection".

[†]Atallah et al., "Detection of Significant Sets of Episodes in Event Sequences"; Gwadera et al., "Reliable detection of episodes in event sequences", "Markov models for identification of significant episodes".

[‡]Tatti, "Significance of Episodes Based on Minimal Windows".

Approaches: kind of episodes

- General episodes have severe pattern explosion: a 9-event serial episode has over 140 million general subepisodes* but only 510 serial subepisodes
- Deciding if a general episode occurs in a sequence is NP-complete, although cases seen in practice are easy to solve†
- Many algorithms mine only parallel or serial episodes, which are easier to handle than general episodes
- However, there are intermediates between the extremes‡

*Tatti and Cule, “Mining closed strict episodes”.

†Tatti and Cule, “Mining closed episodes with simultaneous events”.

‡E.g. the above and Fournier-Viger et al., “Mining Partially-Ordered Episode Rules in an Event Sequence”

The WINEPI algorithm

- WINEPI* is the 'original' episode mining algorithm
- Mines frequent episodes and their rules according to the window-based definition of frequency
- Variations for parallel and serial, injective and non-injective episodes

*Mannila et al., "Discovering Frequent Episodes in Sequences", "Discovery of frequent episodes in event sequences".

The WINEPI algorithm

Goal: find the set of frequent episodes

$$\mathcal{F} = \{A \mid \mathbb{P}(A) \geq \mathbb{P}_{\text{thresh}}\},$$

and the set of episode rules

$$\mathcal{R} = \{B \rightarrow A \mid A \in \mathcal{F}, B \preceq A, \varphi(B \rightarrow A) \geq \varphi_{\text{thresh}}\},$$

given the kind of episodes to mine, the event sequence s , the window size win , the frequency threshold $\mathbb{P}_{\text{thresh}}$, and the confidence threshold φ_{thresh} .

The WINEPI algorithm

$\mathcal{C}_1 \leftarrow \{\{L(e)\} \mid e \in \mathbf{s}\}$ \triangleright Base case of candidate generation

$l \leftarrow 1$

while $\mathcal{C}_l \neq \emptyset$ **do**

$\mathcal{F}_l \leftarrow \text{Recognise}(\mathcal{C}_l, \mathbf{s}, \text{win}, \mathbb{P}_{\text{thresh}})$

$\mathcal{C}_{l+1} \leftarrow \text{GenerateCandidates}(\mathcal{F}_l)$

$l += 1$

Output $\mathcal{F} = \bigcup_{l=1}^{\infty} \mathcal{F}_l$

Output $\mathcal{R} = \text{GenerateRules}(\mathcal{F}, \varphi_{\text{thresh}})$

Does this algorithm look familiar?

The WINEPI algorithm

$\mathcal{C}_1 \leftarrow \{\{L(e)\} \mid e \in \mathbf{s}\}$ \triangleright Base case of candidate generation

$l \leftarrow 1$

while $\mathcal{C}_l \neq \emptyset$ **do**

$\mathcal{F}_l \leftarrow \text{Recognise}(\mathcal{C}_l, \mathbf{s}, \text{win}, \mathbb{P}_{\text{thresh}})$

$\mathcal{C}_{l+1} \leftarrow \text{GenerateCandidates}(\mathcal{F}_l)$

$l += 1$

Output $\mathcal{F} = \bigcup_{l=1}^{\infty} \mathcal{F}_l$

Output $\mathcal{R} = \text{GenerateRules}(\mathcal{F}, \varphi_{\text{thresh}})$

Does this algorithm look familiar? It's APRIORI!

WINEPI: recognition pass

Goal: given \mathcal{C}_l , the set of complexity l candidate episodes, output \mathcal{F}_l , the set of complexity l frequent episodes and their frequencies:


$$\{(A, \mathbb{P}(A)) \in \mathcal{C}_l \times [\mathbb{P}_{\text{thresh}}, 1]\}$$

Sliding window approach: as the window slides by one unit, we

- process events that just entered the window: keep track of which episodes are now inside the window
- process events that just exited the window: keep track of which episodes are no longer inside the window and update their frequencies

WINEPI: recognition pass


Consider the parallel episode $A = \{a, b, c\}$
 A does not occur: event type c is missing



WINEPI: recognition pass


Consider the parallel episode $A = \{a, b, c\}$

A occurs: event $(c, 5)$ entered




WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$
 A occurs: event $(b, 2)$ exited but $(b, 4)$ remains



WINEPI: recognition pass


Consider the parallel episode $A = \{a, b, c\}$
 A no longer occurs: event $(a, 3)$ exited



WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A




There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

Consider the parallel episode $A = \{a, b, c\}$

Two distinct windows contain A



There are eight windows in total, so $\mathbb{P}(A) = \frac{2}{8}$

WINEPI: recognition pass

How we keep track of episodes depends on if we are mining parallel or serial ones


Parallel: for each parallel episode, count the number of events that are inside the window and have event types contained in the episode

- Once all event types' counts reach the required multiplicity, the episode occurs
- Once any event type's count falls below the required multiplicity, the episode no longer occurs

WINEPI: recognition pass

How we keep track of episodes depends on if we are mining parallel or serial ones

Serial: we identify each serial episode with a non-deterministic finite state machine



FSM for episode $A = (b, a, b, c)$

- When an instance of the FSM is in the final state, the episode occurs
- An FSM instance is deinitialised w_{in} time steps after moving out of the initial state

WINEPI: candidate generation


Goal: given \mathcal{F}_l , the set of complexity l frequent episodes,
output \mathcal{C}_{l+1} , the set of candidate episodes of complexity $l + 1$

```
for all  $\mathcal{B} \in \text{Blocks}(\mathcal{F}_l)$  do    ▷ Episodes sharing first  $l - 1$  events
  for all  $(I, J) \in \text{Pairs}(\mathcal{B})$  do
     $A \leftarrow I \oplus \text{Last}(J)$     ▷ Combine two episodes
    if  $\text{DirectSubepisodes}(A) \subseteq \mathcal{F}_l$  then    ▷ Monotonicity
       $\mathcal{C}_{l+1} \leftarrow \mathcal{C}_{l+1} \cup \{A\}$ 
return  $\mathcal{C}_{l+1}$ 
```


■ If mining parallel episodes, give event types a fixed order

$\text{Pairs}(\mathcal{B})$	Parallel	Serial
Non-injective	2-Combinations \mathcal{B} w/ replacement	$\mathcal{B} \times \mathcal{B}$
Injective	2-Combinations \mathcal{B} w/out replacement	2-Permutations


WINEPI: candidate generation



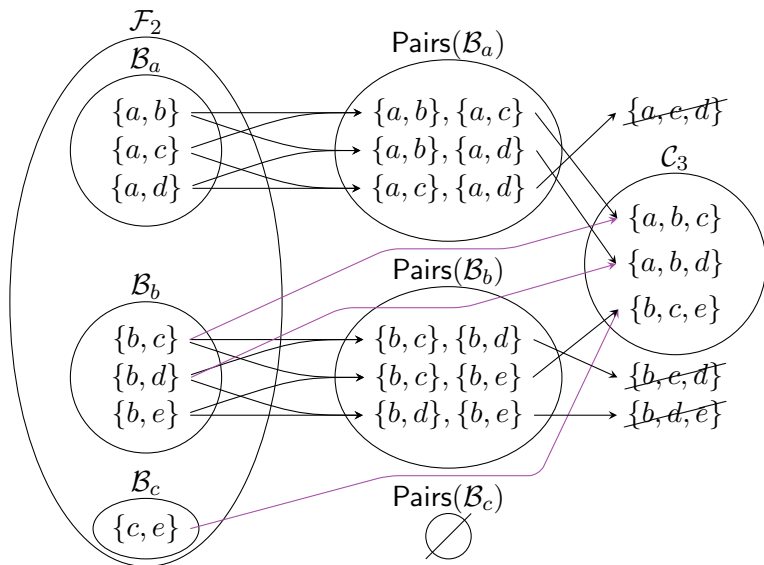
WINEPI: candidate generation



WINEPI: candidate generation



WINEPI: candidate generation



WINEPI: rule generation

Goal: given \mathcal{F} , the set of frequent episodes, and a confidence threshold φ_{thresh} , output \mathcal{R} , the rules that hold in \mathcal{F} w.r.t. φ_{thresh}

```
for all  $A \in \mathcal{F}$  such that  $|A| > 1$  do  
  for all  $B \in \text{DirectSubepisodes}(A)$  do  
    if  $\varphi = \frac{\mathbb{P}(A)}{\mathbb{P}(B)} \geq \varphi_{\text{thresh}}$  then  
       $\mathcal{R} \leftarrow \mathcal{R} \cup \{\text{Rule}(B \rightarrow A, \text{confidence } \varphi)\}$   
return  $\mathcal{R}$ 
```

WINEPI: rule generation

Episode A	$\mathbb{P}(A)$	Rule $B \rightarrow A$	$\varphi(B \rightarrow A)$
$\{a\}$	0.624	$\{a\} \rightarrow \{a, b\}$	$0.029/0.624 = 0.046$
$\{b\}$	0.031	$\{b\} \rightarrow \{a, b\}$	$0.029/0.031 = 0.935$
$\{c\}$	0.167	$\{a\} \rightarrow \{a, c\}$	$0.100/0.624 = 0.160$
$\{a, b\}$	0.029	$\{c\} \rightarrow \{a, c\}$	$0.100/0.167 = 0.599$
$\{a, c\}$	0.100	$\{b\} \rightarrow \{b, c\}$	$0.007/0.031 = 0.226$
$\{b, c\}$	0.007	$\{c\} \rightarrow \{b, c\}$	$0.007/0.167 = 0.042$
$\{a, b, c\}$	0.004	$\{a, b\} \rightarrow \{a, b, c\}$	$0.004/0.029 = 0.138$
		$\{a, c\} \rightarrow \{a, b, c\}$	$0.004/0.100 = 0.040$
		$\{b, c\} \rightarrow \{a, b, c\}$	$0.004/0.007 = 0.571$

The problem with confidence

Episode A	$\mathbb{P}(A)$	Rule $B \rightarrow A$	$\varphi(B \rightarrow A)$
$\{a\}$	0.624	$\{a\} \rightarrow \{a, b\}$	$0.029/0.624 = 0.046$
$\{b\}$	0.031	$\{b\} \rightarrow \{a, b\}$	$0.029/0.031 = 0.935$
$\{c\}$	0.167	$\{a\} \rightarrow \{a, c\}$	$0.100/0.624 = 0.160$
$\{a, b\}$	0.029	$\{c\} \rightarrow \{a, c\}$	$0.100/0.167 = 0.599$
$\{a, c\}$	0.100	$\{b\} \rightarrow \{b, c\}$	$0.007/0.031 = 0.226$
$\{b, c\}$	0.007	$\{c\} \rightarrow \{b, c\}$	$0.007/0.167 = 0.042$
$\{a, b, c\}$	0.004	$\{a, b\} \rightarrow \{a, b, c\}$	$0.004/0.029 = 0.138$
		$\{a, c\} \rightarrow \{a, b, c\}$	$0.004/0.100 = 0.040$
		$\{b, c\} \rightarrow \{a, b, c\}$	$0.004/0.007 = 0.571$

The problem with confidence

- Rule $\{c\} \rightarrow \{a, c\}$ has the second-highest confidence, 0.599
- So is it a good rule?

The problem with confidence

- Rule $\{c\} \rightarrow \{a, c\}$ has the second-highest confidence, 0.599
- So is it a good rule? **No!**
- Confidence has the same problem with episode rules as with association rules
- $\mathbb{P}(\{a, c\}|\{c\}) = 0.599 < 0.624 = \mathbb{P}(\{a\})$

Other goodness measures

- The definitions of goodness measures, such as lift γ , leverage δ , mutual information MI , and Fisher's exact test p_F can be straightforwardly extended to parallel injective episodes*
- Other kinds of episodes require more care: how to compute expected probability of a occurring after b assuming statistical independence?

*Rinta-Paavola, "Identifying Interesting Episode Patterns in User Interaction Log Data".

Redundant episode rules

- An episode rule is redundant if it conveys no more information than a simpler, more general rule
- Let A, B be disjoint parallel episodes and $D \preceq A$. The specialised rule $D \cup B \rightarrow A \cup B$ is **superfluous** given the more general rule $D \rightarrow A$ if the improvement on goodness measures is not statistically significant
- Test using conditional versions of Fisher's exact test or G -test $p_F(B \rightarrow B \cup C \mid D)$, $G = 2 \cdot \text{fr}(D) \cdot MI(B \rightarrow B \cup C \mid D)$ where $C = A \setminus D$
- Note: much of the theory for dependency rules* applies also to parallel injective episode rules

*see e.g. Hämmäläinen and Webb, "A tutorial on statistically sound pattern discovery"

Redundant episode rules

- Many episodes can be involved only in superfluous rules, and can be pruned already in the candidate generation pass
- Does not affect the results—we still have to prune the remaining superfluous rules at the end—but can speed up mining by about an order of magnitude*

*Rinta-Paavola, “Identifying Interesting Episode Patterns in User Interaction Log Data”.

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?
If there is $a \in A$ such that $\mathbb{P}(a \mid A \setminus \{a\}) = 1$, then

- all rules with A in the antecedent are superfluous given the corresponding rule with $A \setminus \{a\}$ in the antecedent,
- but it (or a superepisode) can still be a good consequent!
- That is, there may exist a rule $A \setminus \{b\} \rightarrow A$ that is **not** superfluous given $A \setminus \{a, b\} \rightarrow A \setminus \{a\}$

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?

$\{x\}$

$\{y\}$

$\{z\}$

$\{x, y\}$

$\{x, z\}$

$\{y, z\}$


$\{x, y, z\}$

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are **all** rules $B \rightarrow A$ be superfluous?



Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?

$\{x\}$

$\{y\}$

$\{z\}$

$\{x, y\}$

$\{x, z\}$

$\{y, z\}$


$\{x, y, z\}$

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



If $\mathbb{P}(z \mid x) = 1$, then $\mathbb{P}(z \mid \{x, y\}) = 1$


Rule $\{x, y\} \rightarrow \{x, y, z\}$ is superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



If $\mathbb{P}(z \mid y) = 1$, then $\mathbb{P}(z \mid \{x, y\}) = 1$


Rule $\{x, y\} \rightarrow \{x, y, z\}$ is superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



If $\mathbb{P}(y \mid x) = 1$, then $\mathbb{P}(y \mid \{x, z\}) = 1$


Rule $\{x, z\} \rightarrow \{x, y, z\}$ is superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



If $\mathbb{P}(y \mid z) = 1$, then $\mathbb{P}(y \mid \{x, z\}) = 1$


Rule $\{x, z\} \rightarrow \{x, y, z\}$ is superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



If $\mathbb{P}(x \mid y) = 1$, then $\mathbb{P}(x \mid \{y, z\}) = 1$


Rule $\{y, z\} \rightarrow \{x, y, z\}$ is superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



If $\mathbb{P}(x \mid z) = 1$, then $\mathbb{P}(x \mid \{y, z\}) = 1$


Rule $\{y, z\} \rightarrow \{x, y, z\}$ is superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?




Informally: when, for all direct subepisodes $B \preceq A$, there is a sub-subepisode that 'makes' the rule $B \rightarrow A$ superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?




Informally: when, for all direct subepisodes $B \preceq A$, there is a sub-subepisode that 'makes' the rule $B \rightarrow A$ superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Assume for sake of example: $A = \{x, y, z\}$

When are all rules $B \rightarrow A$ be superfluous?



Informally: when, for all direct subepisodes $B \preceq A$, there is a sub-subepisode that 'makes' the rule $B \rightarrow A$ superfluous

Redundant episode rules

How do we know an episode A is involved only in superfluous rules?

Given a parallel injective episode A , if

$$\forall a \in A \exists b \in A \setminus \{a\} : \mathbb{P}(A \setminus \{a, b\}) = \mathbb{P}(A \setminus \{a\}) \vee \mathbb{P}(A \setminus \{a, b\}) = \mathbb{P}(A \setminus \{b\}),$$

then all rules involving A are superfluous given the corresponding rule involving $A \setminus \{e\}$ for some $e \in A$.*

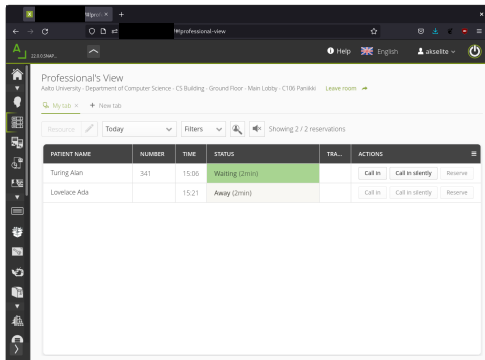
*Rinta-Paavola, "Identifying Interesting Episode Patterns in User Interaction Log Data".

A practical application: my thesis

- For my Master's thesis project*, I used WINEPI enhanced with goodness measures and superfluosness pruning to find interesting patterns of user interaction
- Background: the company, Axel Health, develops a patient flow management system whose purpose is to help patients and healthcare professionals get to the right place at the right time.
- Problem: we have little insight into how users use the system in reality
- Solution: collect log data from the *Professional's View* and analyse it using episode mining methods

*Rinta-Paavola, "Identifying Interesting Episode Patterns in User Interaction Log Data".

The event sequence: simplified example



Professional's View
Aalto University - Department of Computer Science - CS Building - Ground Floor - Main Lobby - C106 Panikki [Leave room](#)

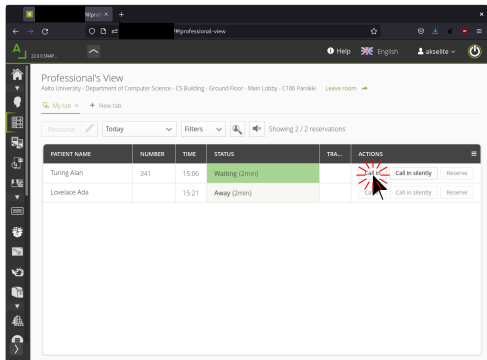
My tab [+ New tab](#)

Resource [Today](#) Filters [Showing 2 / 2 reservations](#)

PATIENT NAME	NUMBER	TIME	STATUS	TRA...	ACTIONS
Turing Alan	341	15:06	Waiting (2min)		Call in Call in silently Reserve
Lovelace Ada		15:21	Away (2min)		Call in Call in silently Reserve

Logged event:
RESERVATION_LIST_SHOW
2023-10-10T17:01:51


The event sequence: simplified example



Professional's View
Aalto University - Department of Computer Science - CS Building - Ground Floor - Main Lobby - C106 Panikki - Leave room

My tab - + New tab

Resource Today Filters Showing 2 / 2 reservations

PATIENT NAME	NUMBER	TIME	STATUS	TRA...	ACTIONS
Turing Aam	341	15:06	Waiting (2min)		 Call in <input type="button" value="Call in silently"/> <input type="button" value="Reserve"/>
Lovelace Ada		15:21	Away (2min)		<input type="button" value="Call in"/> <input type="button" value="Call in silently"/> <input type="button" value="Reserve"/>

Logged event:
CALL_IN(IS_SILENT=false)
2023-10-10T17:15:07

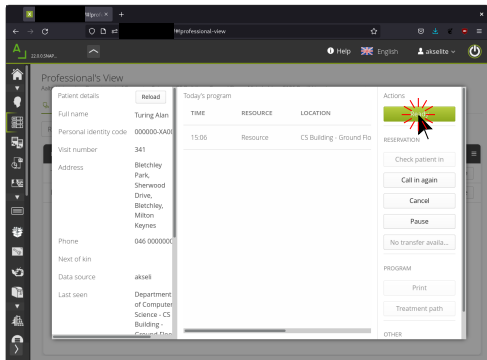
The event sequence: simplified example

The screenshot shows a web application titled "Professional's View" with a dark header bar. The main content area is divided into three sections:

- Patient details:** A form with fields for Full name (Turing Alan), Personal identity code (000000-XAD0), Visit number (341), Address (Bletchley Park, Sherwood Drive, Bletchley, Milton Keynes), Phone (046 0000000), Next of kin, Data source (aksell), and Last seen (Department of Computer Science - CS Building - Ground Floor).
- Today's program:** A table with columns TIME, RESOURCE, and LOCATION. It contains one row: 15:06, Resource, CS Building - Ground Floor.
- Actions:** A vertical list of buttons: Ready (highlighted in green), Check patient in, Call in again, Cancel, Pause, No transfer availa..., Print, Treatment path, and OTHER.

Logged event:
RESERVATION_WINDOW
2023-10-10T17:15:08

The event sequence: simplified example



Logged event:
CALL_IN_CLOSE
2023-10-10T17:59:27

Experiments

- Collected six weeks of data from six healthcare organisers: in total, more than two million events from more than three thousand users
- Preprocessing: parametric events were turned into two, one with parameters and one without; 'uninteresting' events discarded; concatenated event sequences from different users
- Mined episodes with frequency threshold $\mathbb{P}_{\text{thresh}} = 0.001$; tried window sizes from one minute to one hour
- Mined rules with confidence threshold $\varphi_{\text{thresh}} = 0$, so that we can analyse other goodness measures
- Post-processing: discarded rules that were technical artefacts of how the system works, that did not show a statistical dependence ($\delta \leq 0$), or that were superfluous at the significance threshold $\alpha = 0.01$

Results

- Top rules were obvious, corresponding to basic functionality
- Digging a bit deeper, we found some practically relevant rules that reveal parts of the UI that users have trouble with
- All goodness measures have their place; Fisher's exact test and G -test gave nearly identical results
- Superfluosness pruning was vital for both performance and ease of interpretation
- Choice of window size can be difficult; more rules found at larger sizes, but the rules in common at two sizes were ranked similarly

What I learned

- Pattern explosion is a difficult problem: despite pruning reducing the number of rules by up to over an order of magnitude, still impossible to go through all results
- Frequency-based mining is not ideal: most frequent patterns can easily be far from the most interesting ones, seldom-used features got drowned out in the noise
- Statistical significance has its problems: with enough data, you will find statistically significant differences that are so small they have no real-life significance*

*Example: one dataset had a rule that improved confidence and lift of its more general rule by 0.03 %, and was worse by other measures, but was deemed non-superfluous with $p \approx 4 \cdot 10^{-971}$

References I



Atallah, M., Gwadera, R., & Szpankowski, W. (2004). Detection of significant sets of episodes in event sequences. In R. Rastogi, K. Morik, M. Bramer & X. Wu (Eds.), *Proceedings of the fourth IEEE international conference on data mining* (pp. 3–10). IEEE Computer Society.
<https://doi.org/10.1109/icdm.2004.10090>



Casas-Garriga, G. (2003). Discovering unbounded episodes in sequential data. In N. Lavrač, D. Gamberger, L. Todorovski & H. Blockee (Eds.), *Proceedings of the 7th European conference on principles and practice of knowledge discovery in databases. Lecture notes in artificial intelligence: Vol 2838* (pp. 83–94). Springer.
https://doi.org/10.1007/978-3-540-39804-2_10

References II



Fournier-Viger, P., Chen, Y., Nouioua, F., & Lin, J. C.-W. (2021). Mining partially-ordered episode rules in an event sequence. In N. T. Nguyen, S. Chittayasothorn, D. Niyato & B. Trawiński (Eds.), *Proceedings of the 13th Asian conference on intelligent information and database systems. Lecture notes in artificial intelligence, Vol 12672* (pp. 3–15). Springer.

https://doi.org/10.1007/978-3-030-73280-6_1



Gwadera, R., Atallah, M., & Szpankowski, W. (2003). Reliable detection of episodes in event sequences. In X. Wu, A. Tuzhilin & J. Shavlik (Eds.), *Proceedings of the third IEEE international conference on data mining* (pp. 67–74). IEEE Computer Society.

<https://doi.org/10.1109/icdm.2003.1250904>

References III



Gwadera, R., Atallah, M., & Szpankowski, W. (2005). Markov models for identification of significant episodes. In H. Kargupta, J. Srivastava, C. Kamath & A. Goodman (Eds.), *Proceedings of the fifth SIAM international conference on data mining* (pp. 404–414). Society for Industrial and Applied Mathematics.
<https://doi.org/10.1137/1.9781611972757.36>



Hämäläinen, W., & Webb, G. I. (2018). A tutorial on statistically sound pattern discovery. *Data Mining and Knowledge Discovery*, 33(2), 325–377.
<https://doi.org/10.1007/s10618-018-0590-x>

References IV



Laxman, S., Sastry, P. S., & Unnikrishnan, K. P. (2005). Discovering frequent episodes and learning hidden Markov models: A formal connection. *IEEE Transactions on Knowledge and Data Engineering*, 17(11), 1505–1517.
<https://doi.org/10.1109/tkde.2005.181>



Mannila, H., Toivonen, H., & Verkamo, A. I. (1995). Discovering frequent episodes in sequences. In U. Fayyad & R. Uthurusamy (Eds.), *Proceedings of the first international conference on knowledge discovery and data mining* (pp. 210–215). The AAAI Press.
<https://cdn.aaai.org/KDD/1995/KDD95-024.pdf>



Mannila, H., Toivonen, H., & Verkamo, A. I. (1997). Discovery of frequent episodes in event sequences. *Data Mining and Knowledge Discovery*, 1(3), 259–289.
<https://doi.org/10.1023/A:1009748302351>

References V



Méger, N., & Rigotti, C. (2004). Constraint-based mining of episode rules and optimal window sizes. In J.-F. Boulicaut, F. Esposito, F. Giannotti & D. Pedreschi (Eds.), *Proceedings of the 8th European conference on principles and practice of knowledge discovery in databases. Lecture notes in artificial intelligence: Vol 3202* (pp. 313–324). Springer. https://doi.org/10.1007/978-3-540-30116-5_30



Rinta-Paavola, J. (2022). *Identifying interesting episode patterns in user interaction log data* [Master's thesis]. Aalto University. School of Science.
<http://urn.fi/URN:NBN:fi:aalto-202301291820>

References VI

-  Tatti, N. (2009). Significance of episodes based on minimal windows. In J.-F. Boulicaut, F. Esposito, F. Giannotti & D. Pedreschi (Eds.), *Proceedings of the ninth IEEE international conference on data mining* (pp. 513–522). IEEE Computer Society.
<https://doi.org/10.1109/icdm.2009.23>
-  Tatti, N., & Cule, B. (2011). Mining closed episodes with simultaneous events. In C. Apte, J. Ghosh & P. Smyth (Eds.), *Proceedings of the 17th ACM SIGKDD international conference on knowledge discovery and data mining* (pp. 1172–1180). Association for Computing Machinery. <https://doi.org/10.1145/2020408.2020589>
-  Tatti, N., & Cule, B. (2012). Mining closed strict episodes. *Data Mining and Knowledge Discovery*, 25(1), 34–66.
<https://doi.org/10.1007/s10618-011-0232-z>