


Mining graph data: Web and recommender systems



1. Introduction to Graph mining
2. Web mining and search
 - Emphasis: **Ranking algorithms**
3. Recommender systems
 - Emphasis: **Collaborative filtering**

image source: Opte project, <https://www.opte.org/the-internet>

1. Introduction to Graph mining

Data may consist of

1. one large graph
 - e.g., internet, social network
2. multiple small graphs
 - e.g., chemical compounds, biological pathways, program control flows, consumer behaviour, ...

Information to mine: interesting substructures, influential nodes, similarities, communities, clusters

Data: one large graph

Internet (Wikipedia hyperlink structure)

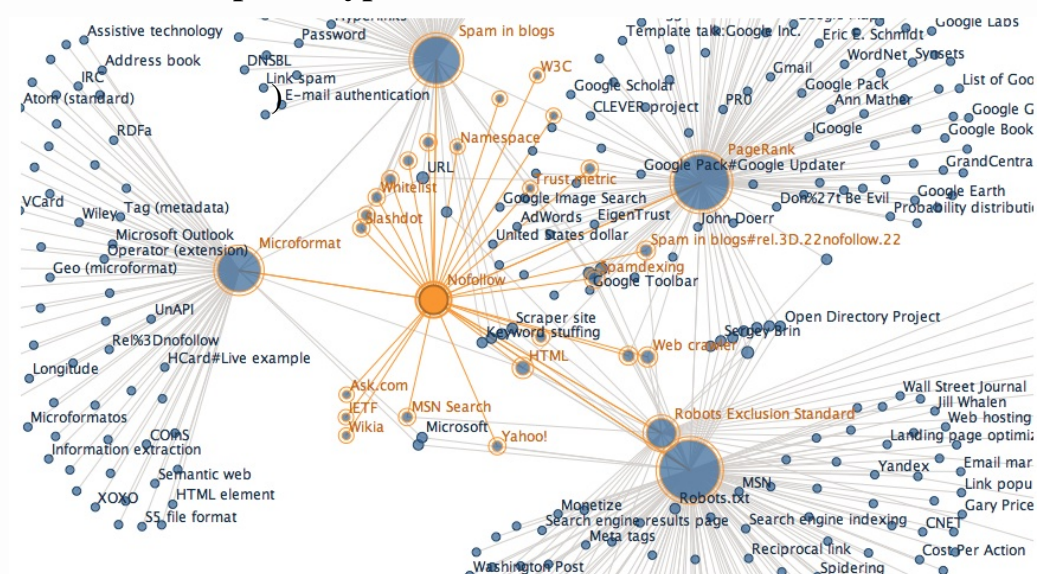


Image source: <https://wiki.digitalmethods.net/Dmi/WikipediaAnalysis>

Social network





image: Lu & Lakshmanan ICDM 2012
<https://www.slideshare.net/WeiLu12/profit-maximization-over-social-networks>

Protein-protein interaction network




Srivastava et al. (2018)
doi 10.2174/1875036201811010240


Most influential nodes?
Communities or dense subgraphs?
Node similarity/future links?

Data: multiple (smaller) graphs

Molecular structure graphs



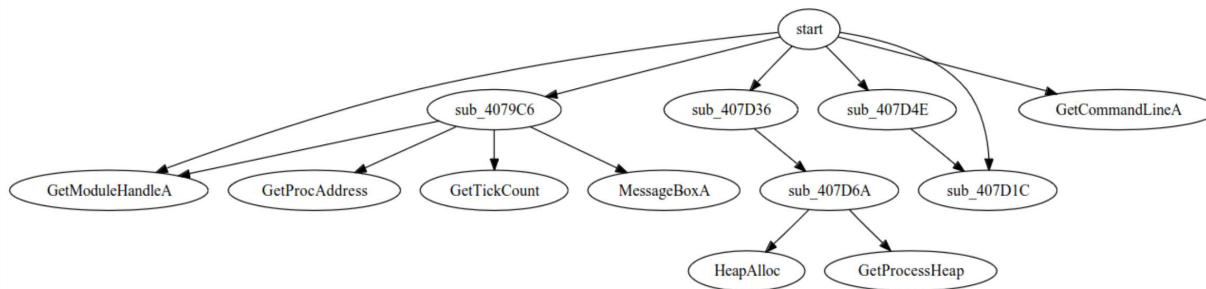
(a) Acetaminophen



(b) Graph representation

Common and unique properties?
Frequent subgraphs?
Clusters?

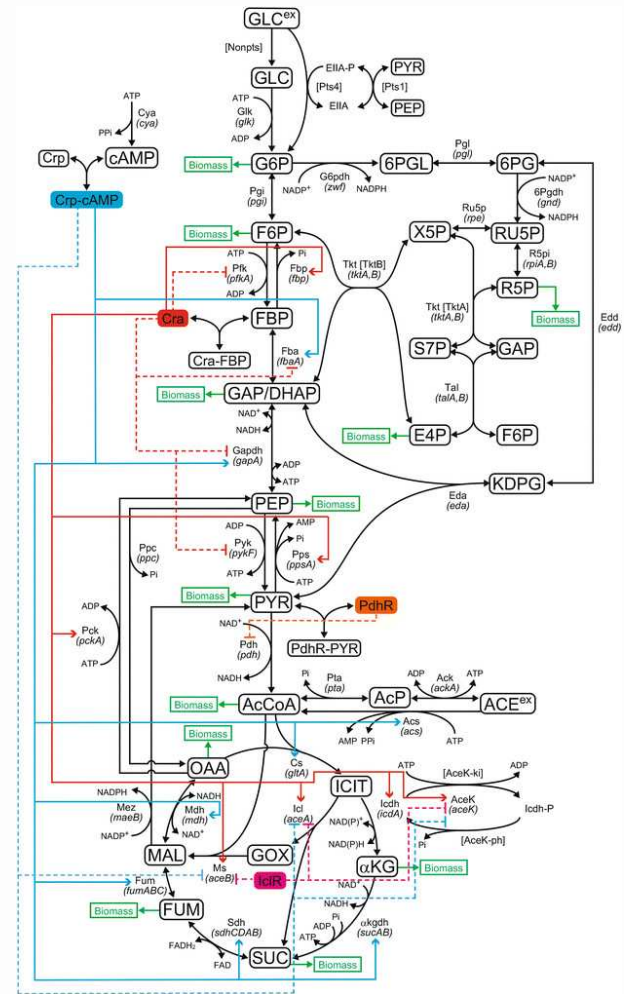
Software call graphs (malware)



Kinable & Kostakis (2011)

doi <https://doi.org/10.1007/s11416-011-0151-y>

Metabolic network (E. coli)




Jahan et al. (2016)

DOI:10.1186/s12934-016-0511-x

Data: graph-form presentation of other data

Similarity graph



User-item utility graph




Image sources MDM2022/Laaksonen and Aggarwal Fig 18.5

2a Web mining

Data:

1. Web content = documents and links
2. Web usage data, like buying and browsing behaviour, ratings, logs

Applications:

1. **Content-centric:** document clustering, resource discovery, web search, linkage mining
2. **Usage-centric:** recommender systems, web log analysis

2b Web search (information retrieval)

1. **Web crawling:** collect all (relevant) documents into a central location (+ keep it up to date) → Aggarwal 18.2
2. **Index construction** for the collection (offline)
3. **Query processing** (online)
 - query = set of keywords $\mathbf{q} = (w_1, \dots, w_k)$
 - identify documents containing all/most $w_i \in \mathbf{q}$
 - **rank** documents (relevance to the query and quality)
 - return best ranked documents

Search engine: steps 2 + 3

Index construction

Preprocess documents, extract relevant tokens (words or phrases) and construct **inverted indices**

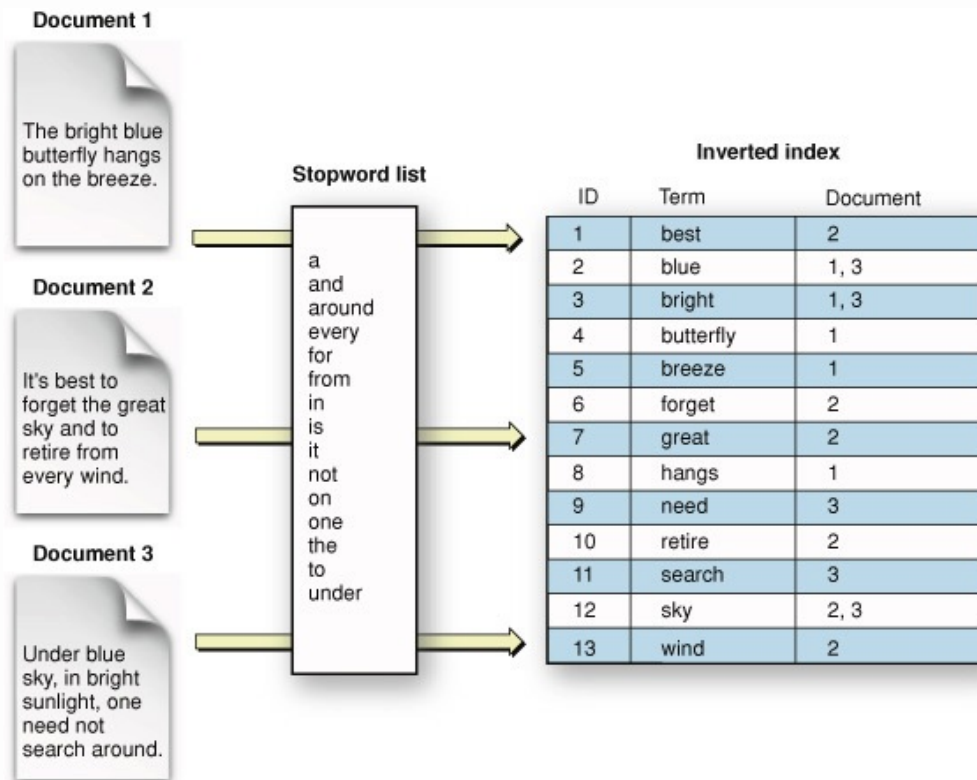


Image source <https://community.hitachivantara.com/s/article/search-the-inverted-index>

Ranking documents

Combine two types of scores:

1. Content-based scores based on occurrence of keywords

- frequency of word
- where the word occurs (more weight if in the title or anchor text)
- prominence of the word (font size, colour)
- relative position of keywords (more relevant if close to each other)

How web spammer can cheat the search engine to get high content-based scores?

Ranking documents

2. Reputation-based scores utilize natural voting mechanisms

a) assumptions based on page citations

- High quality pages are pointed by many pages or
- High quality pages are pointed by many high quality pages

b) user feedback

- users choose relevant pages
- return other similar pages or pages accessed by similar users (recommender systems)

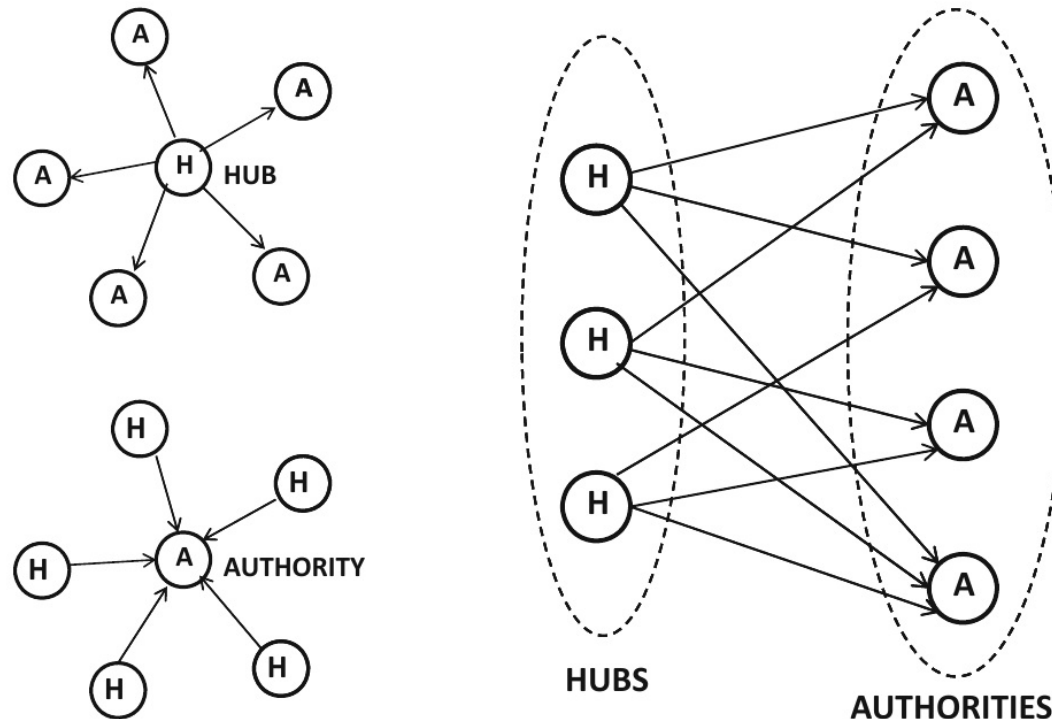
Two ranking algorithms

1. HITS = Hyperlink-Induced Topic Search = "Hubs and Authorities algorithm"
 - query-dependent
 - There are two types of good pages: good authority pages are reputable sources on topics and good hub pages offer links to good sources.
2. PageRank
 - query-independent
 - Highly reputable pages are likely cited by other highly reputable pages.

Hyperlink-Induced Topic Search (HITS)

hub = page that refers to many authoritative pages

authority = authoritative page referred by many hubs



(a) Hub and authority examples

(b) Network organization between hubs and authorities

HITS basic idea

1. Construct an input graph:

- search top- K pages most relevant to the query (e.g., top-200) \rightarrow root set \mathbf{R}
- add pages pointed by $v \in \mathbf{R}$ and some of pages pointing to $v \in \mathbf{R}$ (e.g., max 50/per node) \rightarrow base set \mathbf{V}
- construct a graph $\mathbf{G} = (\mathbf{V}, \mathbf{E})$ where \mathbf{E} contains all links between nodes in \mathbf{V}

2. Assign all nodes hub and authority weights, $h(v)$ and $a(v)$

- update weights until the system converges
- return nodes v with highest $a(v)$


HITS algorithm on $G = (V, E)$

$In(v) = \{u \mid (u, v) \in E\}$ nodes pointing to v

$Out(v) = \{u \mid (v, u) \in E\}$ nodes to which v points

1. initialize h and a such that $\sum_i h(v_i)^2 = \sum_i a(v_i)^2 = 1$
(e.g., $h(v_i) = a(v_i) = \frac{1}{\sqrt{n}}$)
2. until h and a values converge, update for all v_i :
 - $h(v_i) = \sum_{w \in Out(v_i)} a(w)$ (authorities pointed by v_i)
 - $a(v_i) = \sum_{w \in In(v_i)} h(w)$ (hubs pointing to v_i)
 - normalize weights: $h(v_i) = \frac{h(v_i)}{H}$ and $a(v_i) = \frac{a(v_i)}{A}$,
where $H = \sqrt{\sum_i h(v_i)^2}$ ja $A = \sqrt{\sum_i a(v_i)^2}$
3. return v_i s with largest $a(v_i)$

Example: calculation of h and a



initialize $h(i) = a(i) = \frac{1}{\sqrt{4}} = \frac{1}{2}$

round 1:

$$h(1) = \frac{3}{2}, a(1) = 0$$

$$h(2) = \frac{2}{2}, a(2) = \frac{2}{2}$$

$$h(3) = \frac{1}{2}, a(3) = \frac{2}{2}$$

$$h(4) = 0, a(4) = \frac{2}{2}$$

normalize by $H = \sqrt{\frac{7}{2}}$ and $A = \sqrt{3}$

$$h(1) = \frac{3}{\sqrt{14}}, a(1) = 0$$

$$h(2) = \frac{2}{\sqrt{14}}, a(2) = \frac{1}{\sqrt{3}}$$

$$h(3) = \frac{1}{\sqrt{14}}, a(3) = \frac{1}{\sqrt{3}}$$

$$h(4) = 0, a(4) = \frac{1}{\sqrt{3}}$$

Example: final result

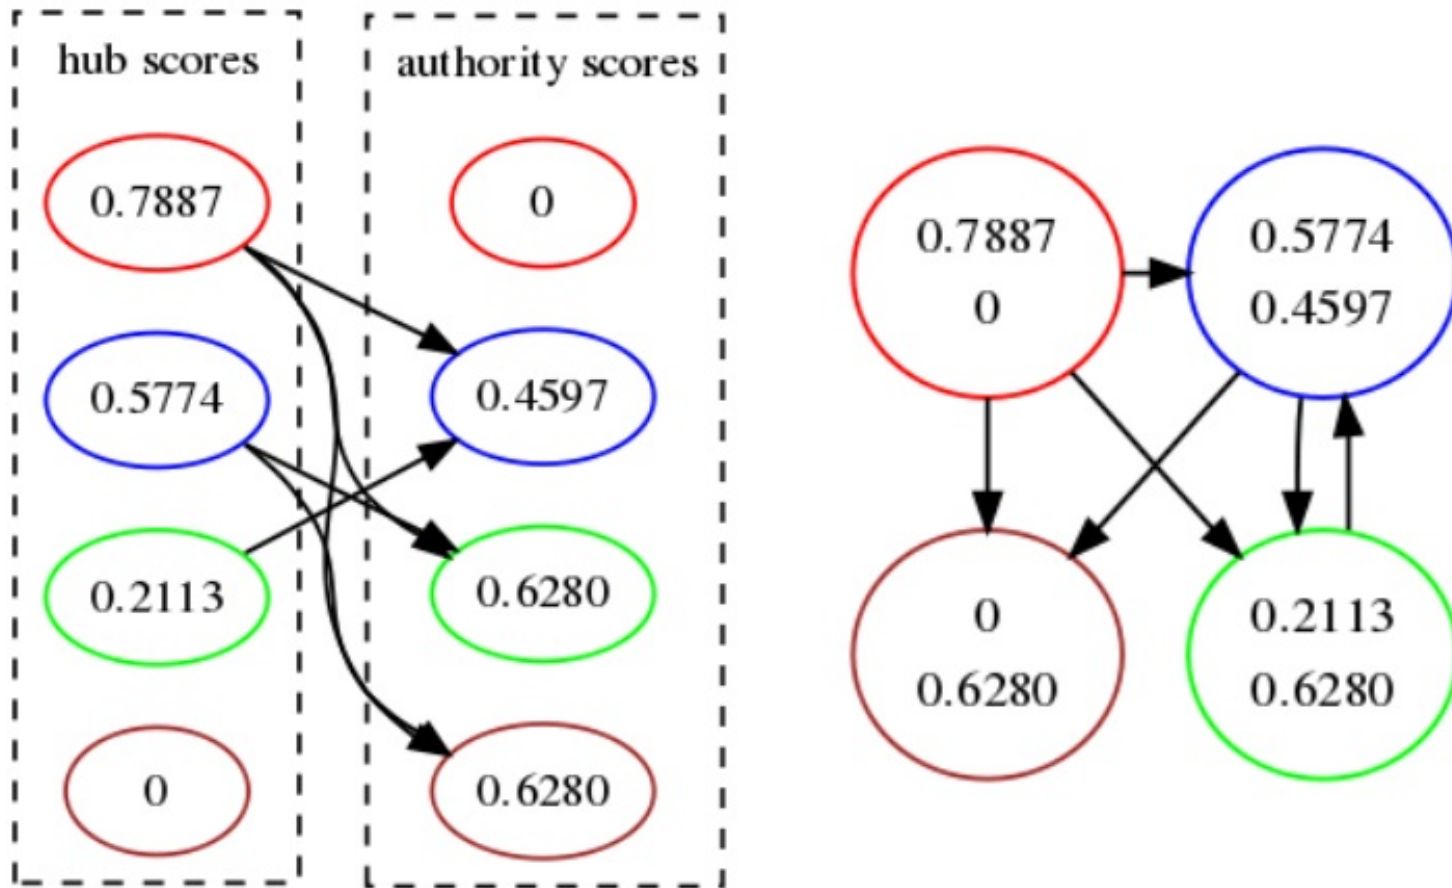


image source Pajarinen (2008): The PageRank/HITS algorithms (slides)

PageRank

Uses a random surfer model = random walk model

- visit random pages by selecting random links
- long term visiting frequency of a page depends on the number of in-linking pages and their visiting frequency
- \Rightarrow frequency is high, if linked from other frequently visited pages
- reputation \approx long term frequency of visits by a random surfer = **steady-state probability** π


\Rightarrow calculate $\pi = (\pi_1, \dots, \pi_n)^T$, where π_i = PageRank of the i th node v_i (n =number of pages)

Problem: the surfer can get trapped!

Originally, all out-going links from a node have equal probability.

Dead-end nodes don't have out-going links

⇒ add links from a dead-end node to all n pages with transition probability $\frac{1}{n}$



Problem: the surfer can get trapped!

Dead-end components don't contain out-going links from the group

⇒ all steady-state probability becomes concentrated into such groups!




image source Aggarwal Fig. 18.2

Teleportation offers a solution to dead-end components

At each transition, a random surfer may

- jump to an arbitrary page with probability α or
- follow one of links with probability $1 - \alpha$

α = **smoothing or damping probability**

- typically $\alpha = 0.1$
- large α makes steady-state probability distribution more uniform
- What happens if $\alpha = 1$?

Presentations as a Markov chain

- graph $\mathbf{G} = (\mathbf{V}, \mathbf{E})$, \mathbf{V} = pages, $|\mathbf{V}| = n$, \mathbf{E} = links
- $In(v) = \{u \mid (u, v) \in \mathbf{E}\}$ nodes pointing to v
- $Out(v) = \{u \mid (v, u) \in \mathbf{E}\}$ nodes pointed by v
- transition matrix \mathbf{P} , where $p_{ij} = \mathbf{P}[i, j]$ probability of transitioning $v_i \rightarrow v_j$
- set $p_{ij} = \frac{1}{|Out(v_i)|}$
- $\pi = (\pi_1, \dots, \pi_n)^T$ steady-state probabilities

$$\pi_i = \frac{\alpha}{n} + (1 - \alpha) \sum_{v_j \in In(v_i)} p_{ji} \cdot \pi_j$$

Computation with the power iteration method

Idea: update π iteratively, notate by $\pi^{(t)}$ the current π at time step t .

Let $\alpha = (\alpha, \alpha, \dots, \alpha)$ (vector of n α s)

- initialize $\pi_i^{(0)} = 1/n, t = 0$
- until $\pi^{(t)}$ converges do
 - i) calculate $\pi^{(t+1)}$:

$$\pi^{(t+1)} = \frac{\alpha}{n} + (1 - \alpha)\mathbf{P}^T \pi^{(t)}$$

- ii) normalize such that $\sum \pi_i = 1$

Computation with the power iteration method

Given convergence threshold θ :

- initialize $\pi_i^{(0)} = 1/n, t = 0$
- until $(\|\pi^{(t)} - \pi^{(t+1)}\| \leq \theta)$ do
 - for all $i = 1, \dots, n$

$$\pi_i^{(t+1)} = \frac{\alpha}{n} + (1 - \alpha) \sum_{v_j \in \text{In}(v_i)} p_{ji} \pi_j^{(t)}$$

- for all $i = 1, \dots, n$ normalize $\pi_i^{(t+1)} \mathbf{s}$

Note: heavy computation! → calculate beforehand for all pages

Alternatively integrate teleportation probabilities into the transition matrix (here $\alpha = 0.1$)




image source

<https://introcs.cs.princeton.edu/python/16pagerank/>

Other PageRank style methods

1. Topic-sensitive PageRank favours certain topics (like user's interests)

- fix a list of base topics
- search high quality sample of pages from each topic
- restrict teleportation only to these pages

2. SimRank measures similarity between nodes

Intuition: Two random surfers walking backwards from nodes i and j take $L(i, j)$ steps, until they meet. Then $\text{SimRank}(i, j)$ is expected value of $c^{L(i, j)}$, where c = decay constant.

SimRank

$$\text{SimRank}(v_i, v_j) =$$

$$\begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } |In(v_i)| = 0 \text{ or } |In(v_j)| = 0 \\ \frac{c}{|In(v_i)| \cdot |In(v_j)|} \sum_{p \in In(v_i)} \sum_{q \in In(v_j)} \text{SimRank}(p, q) & \text{otherwise} \end{cases}$$

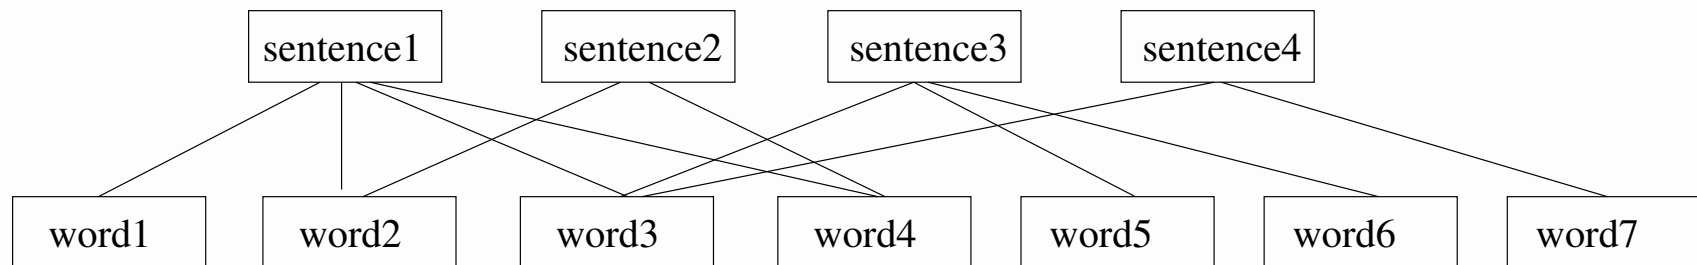
Further reading Aggarwal Ch 18.4.1.2

Summary

- Important problem how to rank documents!
- Content-based scores based on page contents
- Reputation-based scores utilize e.g., link structure
- Ranking algorithms:
 - HITS: query-dependent → smaller input graph
 - PageRank: query-independent → calculate for all pages (heavy!)

Extra: Applying HITS to sentences and words

Given a keyword or words, create a bipartite graph:



- e.g., all sentences where the given word occurs, all their words, all their sentences and randomly some of their words
- sentences (s) are given S weights and words (w) W weights
- initialize uniformly such that the square sum is 1

Extra: Applying HITS to sentences and words

- update rule: $S(s) = \sum_{w_i \in s} W(w_i)$ and $W(w) = \sum_{s_i \in w} S(s_i) +$ normalization
- in the end the words with largest S -weight are most similar
- analogously search the most similar sentences with a given one
- quality of results varies a lot depending on the documents, language, preprocessing, and how the graph was constructed!