


Mining text data

1. Overview
2. Preprocessing text data
3. Representations (bag-of-words)
4. Text clustering and its applications
5. Extra: Word embeddings



1. Overview: Text data

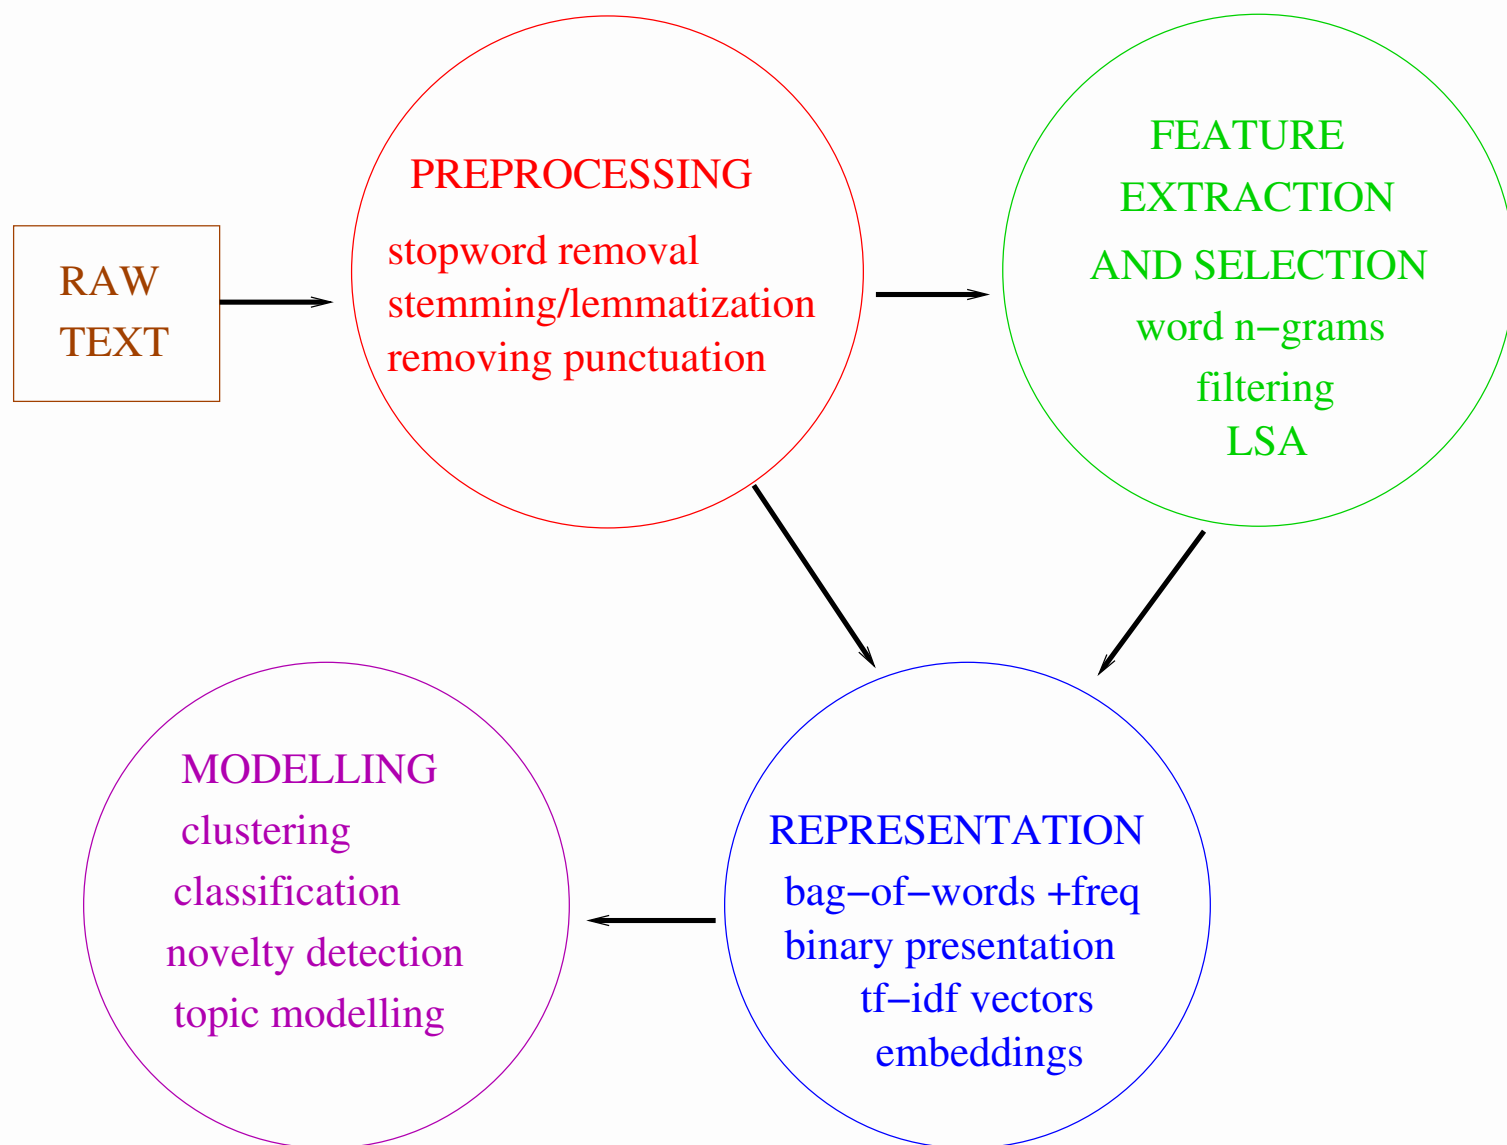
Corpus = collection of text documents $(\mathbf{d}_1, \dots, \mathbf{d}_n)$

Lexicon = set of words (w_1, \dots, w_m)

Document

- **originally ordered**: sequence of words (string), e.g., $\mathbf{d}_1 = (w_{11}, w_{12}, \dots, w_{1q})$, length q
- **bag-of-words model**: **unordered** set of words (+ frequencies), e.g., *Our cat likes the neighbour's cat.* $\rightarrow \{our: 1, cat: 2, likes: 1, the: 1, neighbour's: 1\}$
- **vector space presentation** as a numerical (or binary) vector $\mathbf{x} = (x_1, \dots, x_m)$
 - very sparse! (many 0s)
 - occurrence of words more important than absence

Main steps and example tasks



2. Main tasks of preprocessing

- **Tokenization** (usually word=token)
- **Lower-casing**
- **Remove stopwords** (may be stemmed or not, may contain apostrophes (*it's, you'd*))
- **Reduce inflected forms into a base form:**
 - **Stemming:** cut suffixes (*running* → *run*)
 - **Lemmatization:** derive dictionary form (*was* → *be*, *mice* → *mouse*)
- **Remove punctuation.** Decide how to handle
 - digits (numbers sometimes informative)
 - dashes in compound words (*cat-food* → *catfood* or *cat food*?)

Stopwords

- frequently occurring words with little information about semantic content
- not discriminative
- articles, prepositions, pronouns, auxiliary verbs, ...
- multiple lists available
- Note: relevance of words depends on the context!
→ check & edit!

.
.
.
cannot
cant
co
computer
con
could
couldnt
cry
de
describe
detail
.
.
.

Warning:
may be
relevant



Stemming and lemmatization

Stemming cuts suffixes → root form

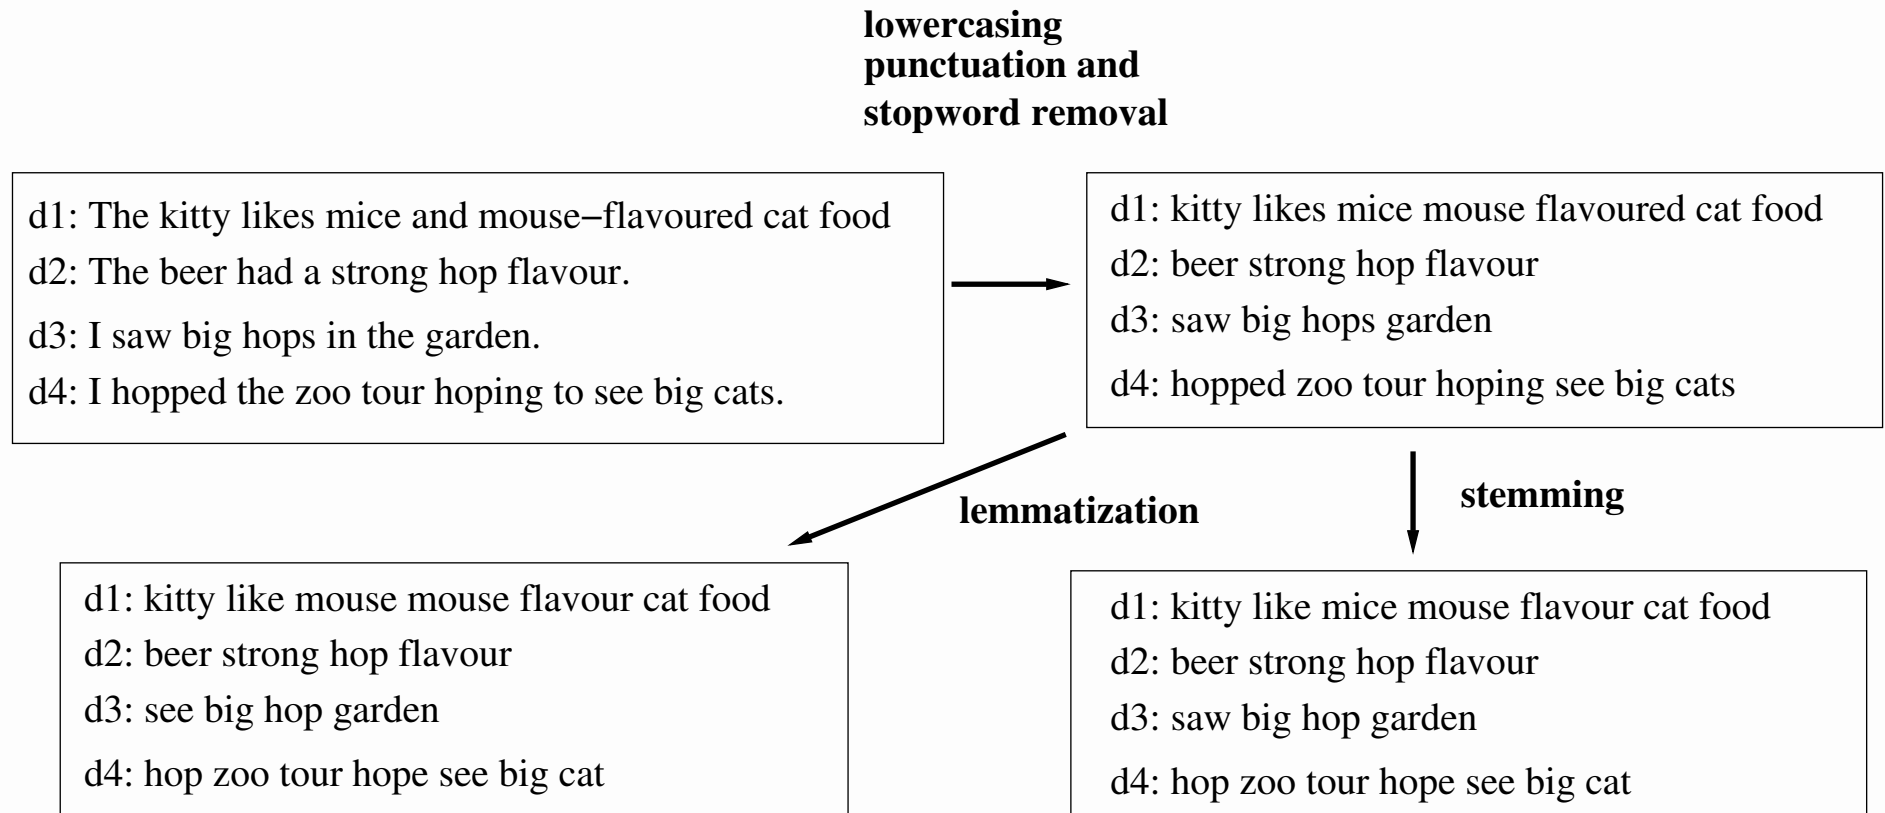
- Overstemming: word is over-truncated → false synonyms (e.g., *hopping* and *hope* → *hop*)
- Understemming: too little truncation → same root not detected (e.g., *alumnus* → *alumnu* vs. *alumni* → *alumni*)
- rule-based, possibly use look-up tables

Lemmatization derive dictionary form

- utilize context, part-of-speech tagging ^a, look-up tables
- more difficult! → slower
- potentially better accuracy

^anoun, verb, adjective, ...; singular/plural; verb tense; subject, object, ...

Example: simple preprocessing



synonymy: kitty, cat

polysemy/homonymy: hop as a verb and a plant



Example: document-term matrix

d1: kitty like mouse mouse flavour cat food

d2: beer strong hop flavour

d3: see big hop garden

d4: hop zoo tour hope see big cat

Typically very sparse matrix!
Present compactly!

| | kitty | like | mouse | flavour | cat | food | beer | strong | hop | see |
|----|-------|------|-------|---------|-----|------|------|--------|-----|-----|
| d1 | 1 | 1 | 2 | 1 | 1 | 1 | | | | |
| d2 | | | | 1 | | | 1 | 1 | 1 | |
| d3 | | | | | | | | | 1 | 1 |
| d4 | | | | | 1 | | | | 1 | 1 |

| | big | garden | zoo | tour | hope |
|----|-----|--------|-----|------|------|
| d1 | | | | | |
| d2 | | | | | |
| d3 | 1 | 1 | | | |
| d4 | 1 | | 1 | 1 | 1 |

Note: 0s often skipped in visual presentation (not stored in real structure)

3. Representation (VSM=vector space model)

1. **Frequency model:** $x_{ij} = fr(w_j|\mathbf{d}_i)$ + possibly normalize to $\|\mathbf{x}\| = 1$ ($fr(w_j|\mathbf{d}_i)$ =number of times w_j occurs in \mathbf{d}_i)
2. **Boolean (binary) model:** $x_{ij} = 1$, if w_j occurs in \mathbf{d}_i and 0 otherwise
3. **Tf-idf representation:** $tf \cdot idf$, where tf =term frequency, idf =inverse document frequency
 - basic form $tfidf(w_j, \mathbf{d}_i) = fr(w_j|\mathbf{d}_i) \cdot \log \frac{n}{n_j}$
 - $df(w_j) = \frac{n_j}{n}$, $idf(w_j) = \frac{n}{n_j}$, n_j =number of documents containing w_j
 - When $tfidf$ is maximal?
 - many variants! → Check what your library calculates!

Some *tf-idf* variants

- *tf* can be normalized: $\frac{fr(w_j|\mathbf{d}_i)}{len(\mathbf{d}_i)}$
- frequency damping: $tf(w_j, \mathbf{d}_i) = \log(fr(w_j|\mathbf{d}_i))$ or $\sqrt{fr(w_j|\mathbf{d}_i)}$
- Note: **sklearn** offers **length-normalized**^a versions of either $fr(w_j|\mathbf{d}_i) \cdot \left(\log\left(\frac{1+n}{1+n_j}\right) + 1\right)$ or $fr(w_j|\mathbf{d}_i) \cdot \left(\log\left(\frac{n}{n_j}\right) + 1\right)$
 - always *idf* > 0
 - $\frac{1+n_j}{1+n}$ smoothed estimate of relative document frequency

^aeither $\sum_j x_{ij}^2 = 1$ or $\sum_j |x_{ij}| = 1$ or none

Example: Basic tf-idf

$tf = fr(w_j | \mathbf{d}_i) = 1 \text{ or } 2$ (*mouse*)

$n_j = df(w_j) \in \{1, 2, 3\} \Rightarrow idf(w_j) \in \{2, 1, 0.415\}$.

| | kitty | like | mouse | flavour | cat | food | beer | strong | hop | see |
|----|-------|------|-------|---------|-----|------|------|--------|-------|-----|
| d1 | 2.0 | 2.0 | 4.0 | 1.0 | 1.0 | 2.0 | | | | |
| d2 | | | | 1.0 | | | 2.0 | 2.0 | 0.415 | |
| d3 | | | | | | | | | 0.415 | 1.0 |
| d4 | | | | | 1.0 | | | | 0.415 | 1.0 |

| | big | garden | zoo | tour | hope |
|----|-----|--------|-----|------|------|
| d1 | | | | | |
| d2 | | | | | |
| d3 | 1.0 | 2.0 | | | |
| d4 | 1.0 | | 2.0 | 2.0 | 2.0 |

Are single words good features?

Basic features: (stemmed) words (excluding stopwords)

1. too frequent words don't separate documents
 - stopword removal and *idf* help
 - common word may be part of useful feature, e.g., *data mining, data management, text data* → word n-grams
2. unique or rare words don't reveal similarity between documents
 - words may still be synonyms (*kitty, cat*) → WordNet
 - or related → LSA and word embeddings can help
 - different spelling (neighbour, neighbor) or errors
3. polysemy problem: similar looking word (or stem) may have different meanings (*hop*)

Create features for word n -grams?

word n -gram = sequence of n words that often occur together (phrases, collocations)

- monogram (unigram, 1-gram): single word “*data*”
- bigram: two words: “*data analysis*”
- trigram: three words: “*Bayesian data analysis*”

results a lot of features! \Rightarrow filtering by

- frequency (keep if $fr(w_1 w_2)$ high)
- significance of association (e.g., $MI(w_1, w_2)$, $\chi^2(w_1, w_2)$)
- dimension reduction (e.g., LSA)

Note: character n -grams = n consecutive characters

Latent semantic analysis (LSA) = SVD applied to text data

- radical dimension reduction (e.g., 100 000 \rightarrow 300)
- helps with synonymy and a bit with polysemy
- similar/related terms tend to be combined, e.g.,
 $\{\text{car, truck, flower}\} \rightarrow \{1.3452 \cdot \text{car} + 0.2828 \cdot \text{truck}, \text{flower}\}$

The diagram illustrates the SVD decomposition of a matrix A (size $n \times d$) into three matrices: U (size $n \times n$), Σ (size $n \times d$), and V^T (size $d \times d$). The matrix A is shown as a pink rectangle. The matrix U is shown as a pink rectangle with a blue rectangle to its right. The matrix Σ is shown as a pink rectangle with a blue rectangle to its right. The matrix V^T is shown as a pink rectangle with a blue rectangle to its right. The resulting approximation \hat{A} is shown as a pink rectangle.

$$\begin{array}{c} A \\ n \times d \end{array} = \begin{array}{c} \hat{U} \\ n \times r \end{array} \begin{array}{c} \hat{\Sigma} \\ r \times r \end{array} \begin{array}{c} \hat{V}^T \\ r \times d \end{array} \rightarrow \hat{A}$$

$U \qquad \Sigma \qquad V^T$
 $n \times n \qquad n \times d \qquad d \times d$

Recall: If old features F_1, \dots, F_d and i th singular vector $(\mathbf{V}_{i1}, \dots, \mathbf{V}_{id})^T$, i th new feature $\sum_{j=1}^d \mathbf{V}_{ij} F_j$. LSA: Recap Aggarwal 2.4.3.3. Image: Perunicic (2017).

4. Clustering text data

Given **vector-space representation** of text data:
any common clustering method

- K -representatives, spectral, hierarchical, probabilistic EM-algorithm, affinity propagation
- distance/similarity: prefer **cosine similarity**!
remember: if data normalized to $\|\mathbf{x}\| = 1$,
 $L_2^2(\mathbf{x}_1, \mathbf{x}_2) = 2(1 - \cos(\mathbf{x}_1, \mathbf{x}_2)) \rightarrow L_2$ ok
- dimension reduction (PCA or LSA) can help

Suggestion: Use K -means always as a baseline to compare other methods!

Example: cos-sim vs. L_2 distance

d1: kitty like mouse mouse flavour cat food

d2: beer strong hop flavour

d3: see big hop garden

d4: hop zoo tour hope see big cat

\cos_T =cos-sim of tfidf vectors, \cos_B =cos-sim of binary vectors,
 L_2 =Euclidean distance between tfidf vectors:

| pair | \cos_T | \cos_B | L_2 | note |
|----------|---------------|--------------|--------------|------------------|
| (d1, d2) | 0.0603 | 0.204 | 6.097 | No common words! |
| (d1, d3) | 0 | 0 | 6.014 | |
| (d1, d4) | 0.0469 | 0.154 | 6.571 | |
| (d2, d3) | 0.0229 | 0.250 | 3.873 | |
| (d2, d4) | 0.0146 | 0.189 | 4.899 | |
| (d3, d4) | 0.2245 | 0.567 | 4.123 | |

Special techniques for clustering text

1. Modifications of K -representatives

- replace cluster centroids by **cluster digests** = most frequent words (e.g., 200–400)
- better strategies to select seeds!

2. Co-clustering

- Idea: cluster words and documents simultaneously
- cluster $C_i = (R_i, V_i)$; V_i = most relevant words (cluster digest) in documents R_i

Note: “cluster digest” has slightly different meanings in different contexts (intention to describe the cluster)

Scatter/Gather: better K -means for text

1. Construct K seeds with agglomerative hierarchical clustering. Two alternatives:
 - a) **Buckshot**: choose \sqrt{Kn} random samples $\rightarrow K$ clusters $\rightarrow K$ seeds
 - b) **Fractionation**: divide data into $\frac{n}{m}$ buckets $\rightarrow \nu m$ clusters per bucket ($\nu \in]0, 1[$) \rightarrow concatenate cluster documents \rightarrow repeat until K clusters left
2. Apply K -means
 - centroid=concatenation of documents + remove infrequent words

Scatter/Gather: better K -means for text

3. (Optionally) **Refine clusters:**

a) Split incoherent clusters

- coherence score: $avg(sim(\mathbf{x}, \mathbf{c}))$ or $avg(sim(\mathbf{x}_i, \mathbf{x}_j) \mid \mathbf{x}_i, \mathbf{x}_j \in \mathbf{C})$
- if too low \rightarrow Buckshot with $K = 2$ + recluster rest

b) Join similar clusters

- topical words overlapping significantly

Cutting et al. (1992). Scatter/Gather: A cluster-based approach to browsing large document collections.

Co-clustering (biclustering)

Idea: rearrange rows and columns of doc-term matrix such that most non-zero entries form blocks.

| | CHAMPION | ELECTRON | TROPHY | RELATIVITY | QUANTUM | TOURNAMENT |
|----------------|----------|----------|--------|------------|---------|------------|
| D ₁ | 2 | 0 | 1 | 1 | 0 | 3 |
| D ₂ | 0 | 2 | 0 | 1 | 3 | 0 |
| D ₃ | 1 | 3 | 0 | 1 | 2 | 0 |
| D ₄ | 2 | 0 | 2 | 0 | 0 | 3 |
| D ₅ | 0 | 2 | 1 | 1 | 3 | 0 |
| D ₆ | 1 | 0 | 2 | 0 | 0 | 3 |

(a) Document-term matrix

| | CHAMPION | TROPHY | TOURNAMENT | ELECTRON | RELATIVITY | QUANTUM |
|----------------|----------|--------|------------|----------|------------|---------|
| D ₁ | 2 | 1 | 3 | 0 | 1 | 0 |
| D ₄ | 2 | 2 | 3 | 0 | 0 | 0 |
| D ₆ | 1 | 2 | 3 | 0 | 0 | 0 |
| D ₂ | 0 | 0 | 0 | 2 | 1 | 3 |
| D ₃ | 1 | 0 | 0 | 3 | 1 | 2 |
| D ₅ | 0 | 1 | 0 | 2 | 1 | 3 |

SPORTS CO-CLUSTER →

← PHYSICS CO-CLUSTER

(b) Re-arranged document-term matrix

=Bipartite graph partitioning problem!

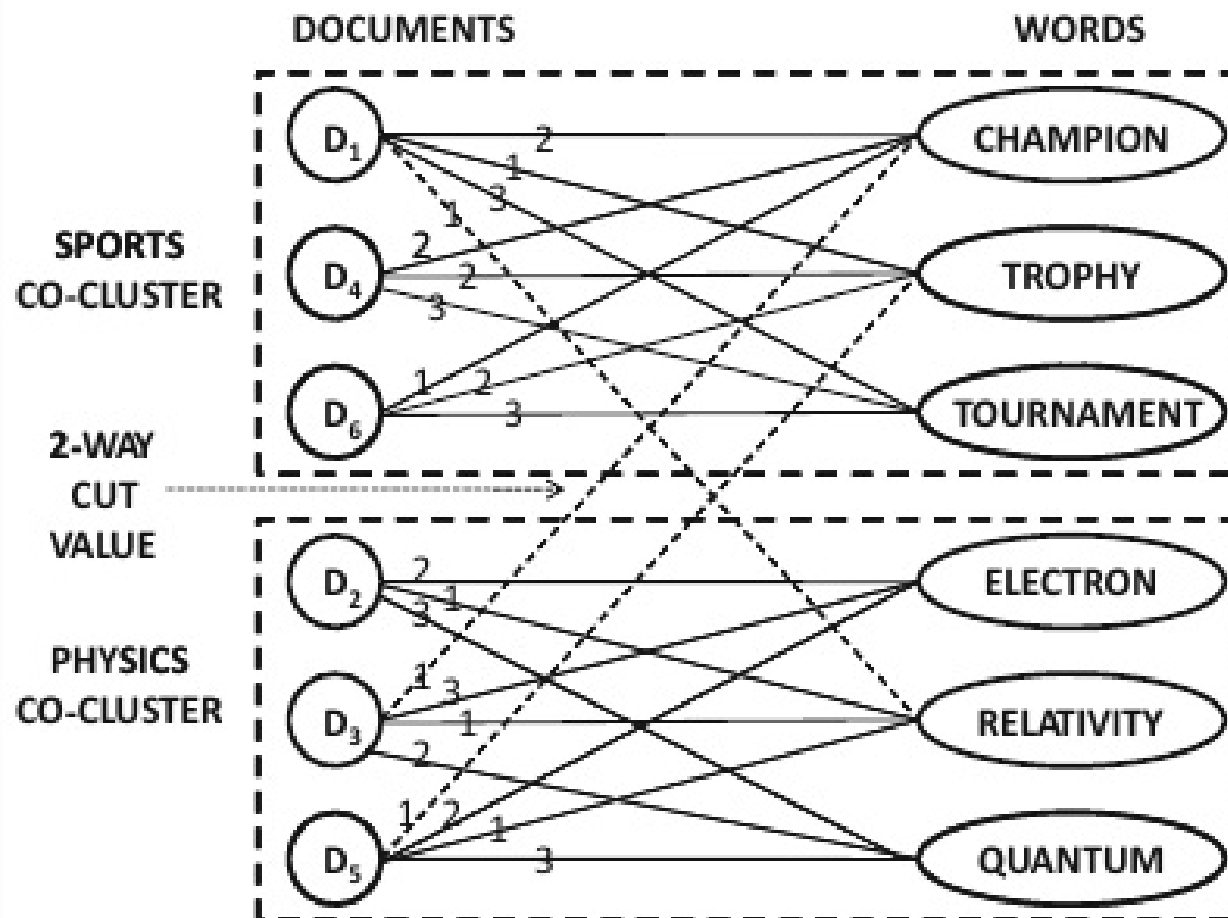


Image source Aggarwal Fig. 13.2

Graph partitioning

Construct graph $\mathbf{G} = (\mathbf{U} \cup \mathbf{V}, \mathbf{E}, \Gamma)$

- \mathbf{U} = nodes for documents (for doc. \mathbf{d}_i node u_i)
- \mathbf{V} = nodes for words (for word w_j node v_j)
- \mathbf{E} = edges
 $(u_i, v_j) \in \mathbf{E}$, if word w_j occurs in \mathbf{d}_i
- Γ = weights
 $\gamma_{ij} = fr(w_j|\mathbf{d}_i)$ or $\gamma_{ij} = tfidf(w_j, \mathbf{d}_i)$

Objective: Partition $\mathbf{U} \cup \mathbf{V}$ into groups such that edge-cut cost ($\sum \gamma_{ij}$ between groups) minimal (+ extra constraints)

Solution: Spectral clustering or other graph partitioning methods (vs. community detection)

Application 1: clustering for classification

Centroid-based classifier:

- Cluster documents of each class (size n_i) into $k_i \propto n_i$ clusters
 - **cluster digest**=most common words of the centroid
 - For \mathbf{d}_{new} :
 - determine K nearest digests (clusters)
 - report the dominant label
-
- + fast alternative to K -NN classifier
 - + handles **synonymy** (similar words \rightarrow same centroid) and **polysemy** (different meanings \rightarrow different centroids)

Application 2: Novelty detection

Problem: Temporal stream of text documents, when a new topic appears? (e.g., news)

Simple solution:

- Maintain a sample of documents, \mathcal{D}
- For \mathbf{d}_{new} calculate $sim(\mathbf{d}_{new}, \mathbf{d})$ for all $\mathbf{d} \in \mathcal{D}$
- If novelty score = $\frac{1}{\max_d sim(\mathbf{d}_{new}, \mathbf{d})}$ high, report \mathbf{d}_{new}

Problem: Pairwise similarity cannot handle synonymy or polysemy \Rightarrow Utilize **micro-clustering**

Novelty detection with micro-clustering

Idea: Maintain K document clusters C_1, \dots, C_K .

For C_i : \mathbf{c}_i = centroid = **cluster digest**, $fr(w_j|C_i)$ word frequencies,
 t_i = time stamp when C_i updated

- given \mathbf{d}_{new} , determine nearest centroid \mathbf{c}_i
- if $(sim(\mathbf{d}_{new}, \mathbf{c}_i) \geq \theta)$
 - add \mathbf{d}_{new} to C_i
 - update $fr(w_j|C_i)$, \mathbf{c}_i (most frequent words) and t_i
- else
 - report \mathbf{d}_{new} as novelty
 - create a new cluster $C = \{\mathbf{d}_{new}\}$ (with new time stamp)
 - remove an old cluster C_i with earliest t_i

5. *Extra: Word embeddings*

Idea: Present words as numerical vectors such that distances reflect semantic similarity

2D presentation of embeddings:






image source Barla (2021)


Word embedding examples



Male-Female



Verb Tense



Country-Capital





image sources Zhang (2020) and Mikolov et al. (2013)

Word embeddings: popular approaches

1. **Matrix factorization:** typically truncated SVD of word co-occurrence (or term-doc) matrix
2. **Word2vec models:** learn a shallow neural network and use its weights
 - a) **CBOW:** predict target word given context words
 - b) **Skip-gram:** predict context words given target word
 - good embeddings require large training sets but models can be transferred to other similar contexts

Mikolov et al. (2013)

Word embeddings: popular approaches

3. **GloVe**: learn embedding vectors \mathbf{x}_i , $\tilde{\mathbf{x}}_i$, such that $\mathbf{x}_i \cdot \tilde{\mathbf{x}}_j \propto \log(P(w_j|w_i))$

- $\tilde{\mathbf{x}}_i$ context embedding \rightarrow later combine with \mathbf{x}_i
- + fast to learn
- + requires less data

Pennington et al. (2014)

Note: *Hype* and *best* are not synonyms. Test always simpler presentations (bag-of-words), too!

Summary

- text presented in vector space (as numerical vectors)
 - simplest: document=bag-of-words with binary occurrence, frequencies, or *tfidf* of words
- preprocessing important → features. Be careful!
- Goal: try to capture important (mid-frequency) words and phrases
 - prune out stopwords, stemming/lemmatization, detect collocations/n-grams, maybe spell-checking
- dimension reduction often useful (+ LSA helps with synonymy/polysemy)
- clustering: K -representatives + modifications, spectral, hierarchical, co-clustering, ...

References & image sources

- Barla (2021): The ultimate guide to word embeddings
<https://neptune.ai/blog/word-embeddings-guide>
- Cutting et al. (1992): Scatter/Gather: A cluster-based approach to browsing large document collections. ACM SIGIR Conference, pp. 318-329.
- Mikolov et al. (2013): Efficient estimation of word representations in vector space. ICLR and arXiv:1301.3781
- Mikolov et al. (2013): Exploiting similarities among languages for machine translation. arXiv:1309.4168

References & image sources

- Pennington et al. (2014): GloVe: Global vectors for word representation. Empirical Methods in Natural Language Processing (EMNLP).
- Perunicic (2017): How are principal component analysis and singular value decomposition related?
<https://intoli.com/blog/pca-and-svd/>
- Zhang (2020): Legal applications of neural word embeddings <https://towardsdatascience.com/legal-applications-of-neural-word-embeddings-556b7515012f>